

Lessons learned from running Terraform at reasonable scale

Why easy, when we can make it complicated?

– the unknown platform engineer

Utilizing FluxCD, Weaveworks TF-Controller and boring-registry at [LYNQTECH](#)

Daniel Ciaglia // Consultant

- **Freelance**
since 2022
- **TIER Mobility SE**
Director of Engineering
- **kreuzwerker GmbH**
Principal Consultant
- **Bundesdruckerei GmbH**
Senior Support Manager
- **[some more]**
- SCUBA dive instructor
- **AWS User Group Berlin**
co-organiser



LinkedIn

Website

Today's menu

1. A typical Terraform stack evolution
2. Running Terraform in GitOps
3. Thoughts on the stack
4. Architectural Decision Records summary

(1.1) Typical Terraform stack evolution¹

Stack: Terraform root module², tracked with 1 state file

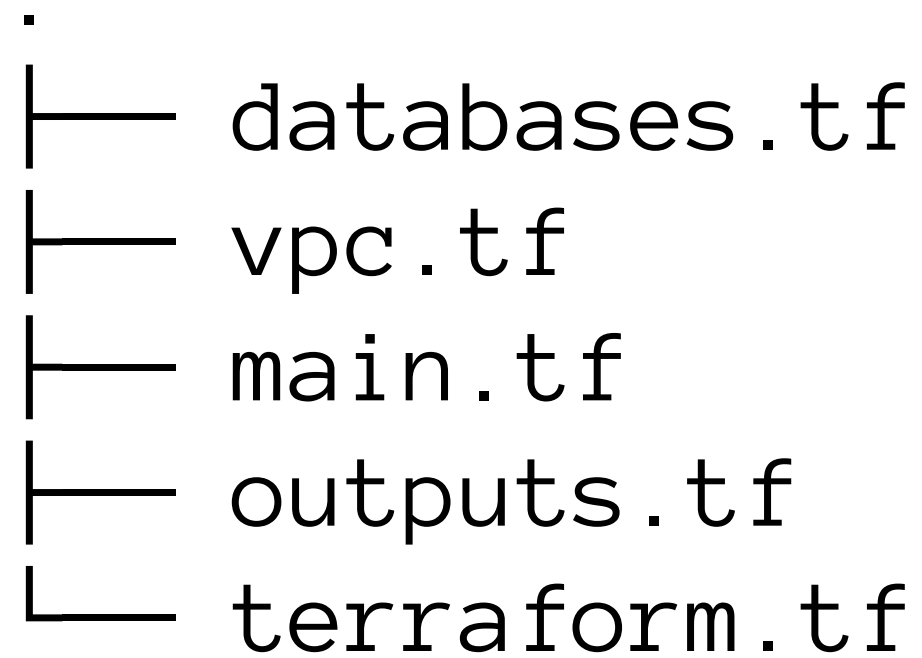
Related: Highly recommend talk "Terraform: from zero to madness" by [@Timur Bublik](#)

¹ your experience might be different 😊

² <https://developer.hashicorp.com/terraform/language/files#the-root-module>

(1.1.1) in the beginning

- you start your project
- put everything in 1 directory
- maybe split files by broader domains.



(1.1.2) The staging/production split

- oh well, you need a staging environment
- both environments are very much the same
- you refactor the code to be parameterised by variables
- you provide 2 .tfvars files

```
.
├── production.tfvars
├── staging.tfvars
├── databases.tf
├── vpc.tf
├── variables.tf
├── main.tf
└── terraform.tf
```

(1.1.3) Code repetition - I need modules

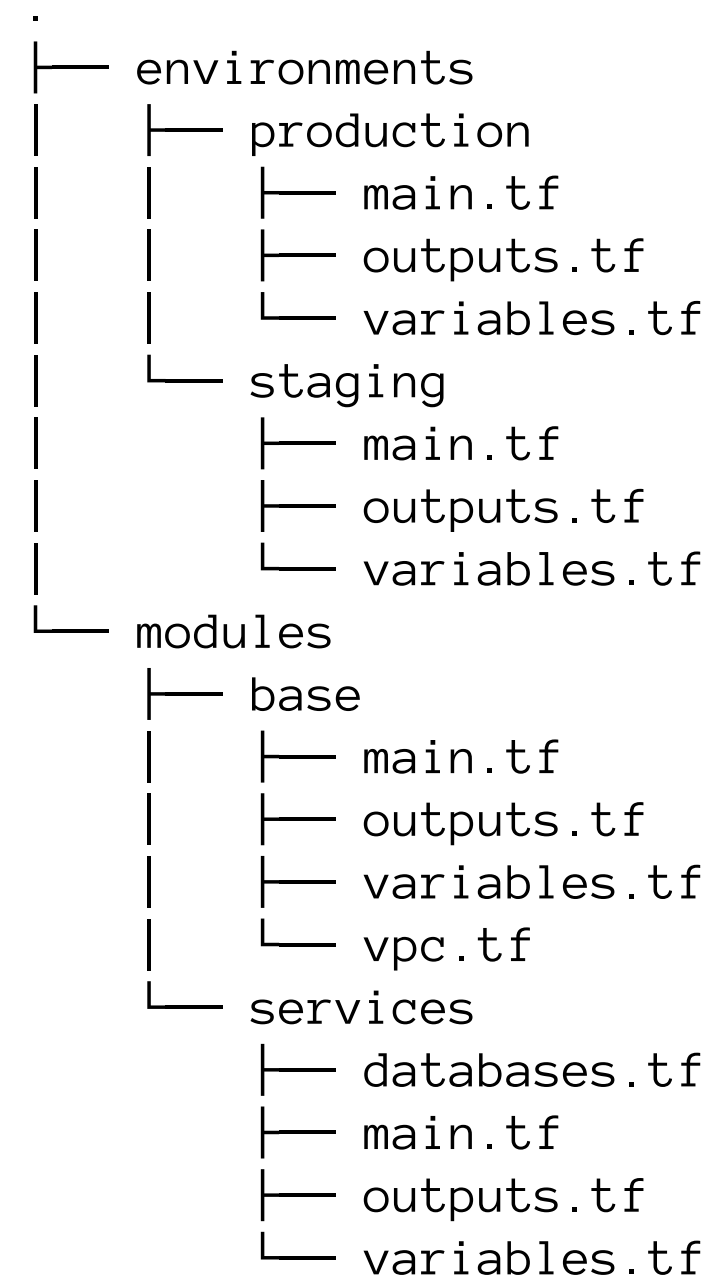
- you add more services and they need infra
- the infra is similar
- you want to keep the code DRY³
- you create a repo, codify best practices, tag them for versioning
- you pull in modules via git

```
# select a specific tag
module "rds" {
  source = "github.com/example/rds?ref=v1.2.0"
}
```

³ <https://en.wikipedia.org/wiki/Don%27trepeatyourself>

(1.1.4) The great separation

- as the stack grows, the environments differ
- you start separating the code in larger blocks
 - the base environment
 - the services
- the code is pulled in as modules
- the services module receives output of base as input
eg. `vpc_id` or `subnets`
- `terraform apply plan` is run manually still



(1.1.5) Fast forward

👉 At this point in time I joined the project 👉

The situation

- as the stack grows further, the amount of resources does as well
- each run of `terraform plan -out plan` takes more and more time
- to review and apply changes for developers becomes a dayfilling job
- you start cheating by targeted apply
- you notice that the amount of files downloaded for each terraform step is enormous⁴
- you notice that git tags can not be used for semantic versioning (version)

Possible solutions

- to address the versioning and data transfer issues - use a private Terraform module registry
- to address the runtime and ownership issue - split the stacks and let the teams handle them (DevOps style)

⁴ HashiTalks DACH 2020 - [Opinionated terraform modules and a registry](#)

(1.2) The boring-registry

- TIER Mobility developed their own "boring" Terraform registry without moving parts (hence the name)
 - Details to be found here: <https://github.com/boring-registry/boring-registry/>
 - The important feature for now is support for the Module Registry Protocol
- **You provide** a S3 bucket, module code and package it in CD via
`./boring-registry upload --type s3 (some more flags) ./your-module`
- **You'll get**
semantic versioning

```
module "rds" {  
  source = "registry.example.com/acme/rds/aws"  
  version = "~> 0.1"  
}
```

(1.3) Separating the service stacks

some Architectural Decisions

Don't

- separate services along team borders⁵
 - ➔ teams and responsibilities change, always
- share states between services
 - ➔ there are secrets in there!⁶
 - ➔ read the docs of the terraform_remote_state data source!

Do

- Layer your stacks - account, network, clusters and services
- 1 Terraform stack per service
 - good for least privilege access
 - place the Terraform code into the service repo
- run the TF stacks in automation
- use an indirect way to share information between stacks⁷

⁵ How TIER switched paradigms - from team- to service-centric

⁶ Sensitive Data in State

⁷ TF-CIX as an approach to share information between terraform stacks

(1.3.1) Indirect information exchange

- use structured data
 - ➔ ideally JSON for `jsondecode()` and `jsonencode()`
- use whatever storage you prefer
 - ➔ SSM Parameter Store or S3

Code for 3 Terraform modules will be provided

- `s3_json_store`
CRUD JSON data on S3
- `ssm_json_store`
CRUD JSON data on SSM Parameter store
- `ssm_json_regex`
read SSM parameter with regex

(1.3.2) Write data (base system)

```
module "ssm_service_data" {
  source  = "registry.example.com/foo/ssm_json_store/aws"
  version = "~> 1.0.2"

  path = "/configuration"
  name = "base"
  data = {
    domain          = local.domain_name
    environment     = local.environment
    environmentClass = local.environmentClass
    backup_plan     = local.backup_plan
    networking = {
      vpc_id          = module.base.vpc_default_id
      subnet_database_ids = module.base.subnet_private_database_ids
      subnet_k8s_ids   = module.base.subnet_private_k8s_ids
    }
  }
  cluster = {
    name          = module.eks.cluster_name
    oidc_issuer_url = module.eks.cluster_oidc_issuer_url
    oidc_provider_arn = module.eks.cluster_oidc_provider_arn
  }
}
```

(1.3.3) Write data (upstream)

```
module "ssm_service_data" {  
  source = "registry.example.com/foo/ssm_json_store/aws"  
  version = "~> 1.0.2"  
  
  path = "/configuration"  
  name = "upstream"  
  data = {  
    installed = true  
  
    private = {}  
    public = {  
      sns = {  
        "foo" = {  
          "arn" = module.sns_foo.arn  
          "name" = module.sns_foo.name  
        }  
      }  
      sqs = {  
        "bar" = {  
          "arn" = module.bar_queue.arn  
          "name" = module.bar_queue.name  
        }  
      }  
    }  
  }  
}
```

(1.3.4) Read data (downstream)

```
module "ssm_data" {
  source = "registry.example.com/foo/ssm_json_store/aws"
  version = "~> 0.1.0"

  path = "/configuration"
  include_filter_regex = "(base|upstream)"
}

module "sns_sqs_subscription_foo" {
  count = try(module.ssm_data.values["upstream"]["installed"], false) ? 1 : 0
  source = "registry.example.com/foo/sns_sqs_subscription/aws"
  version = "~> 0.1"

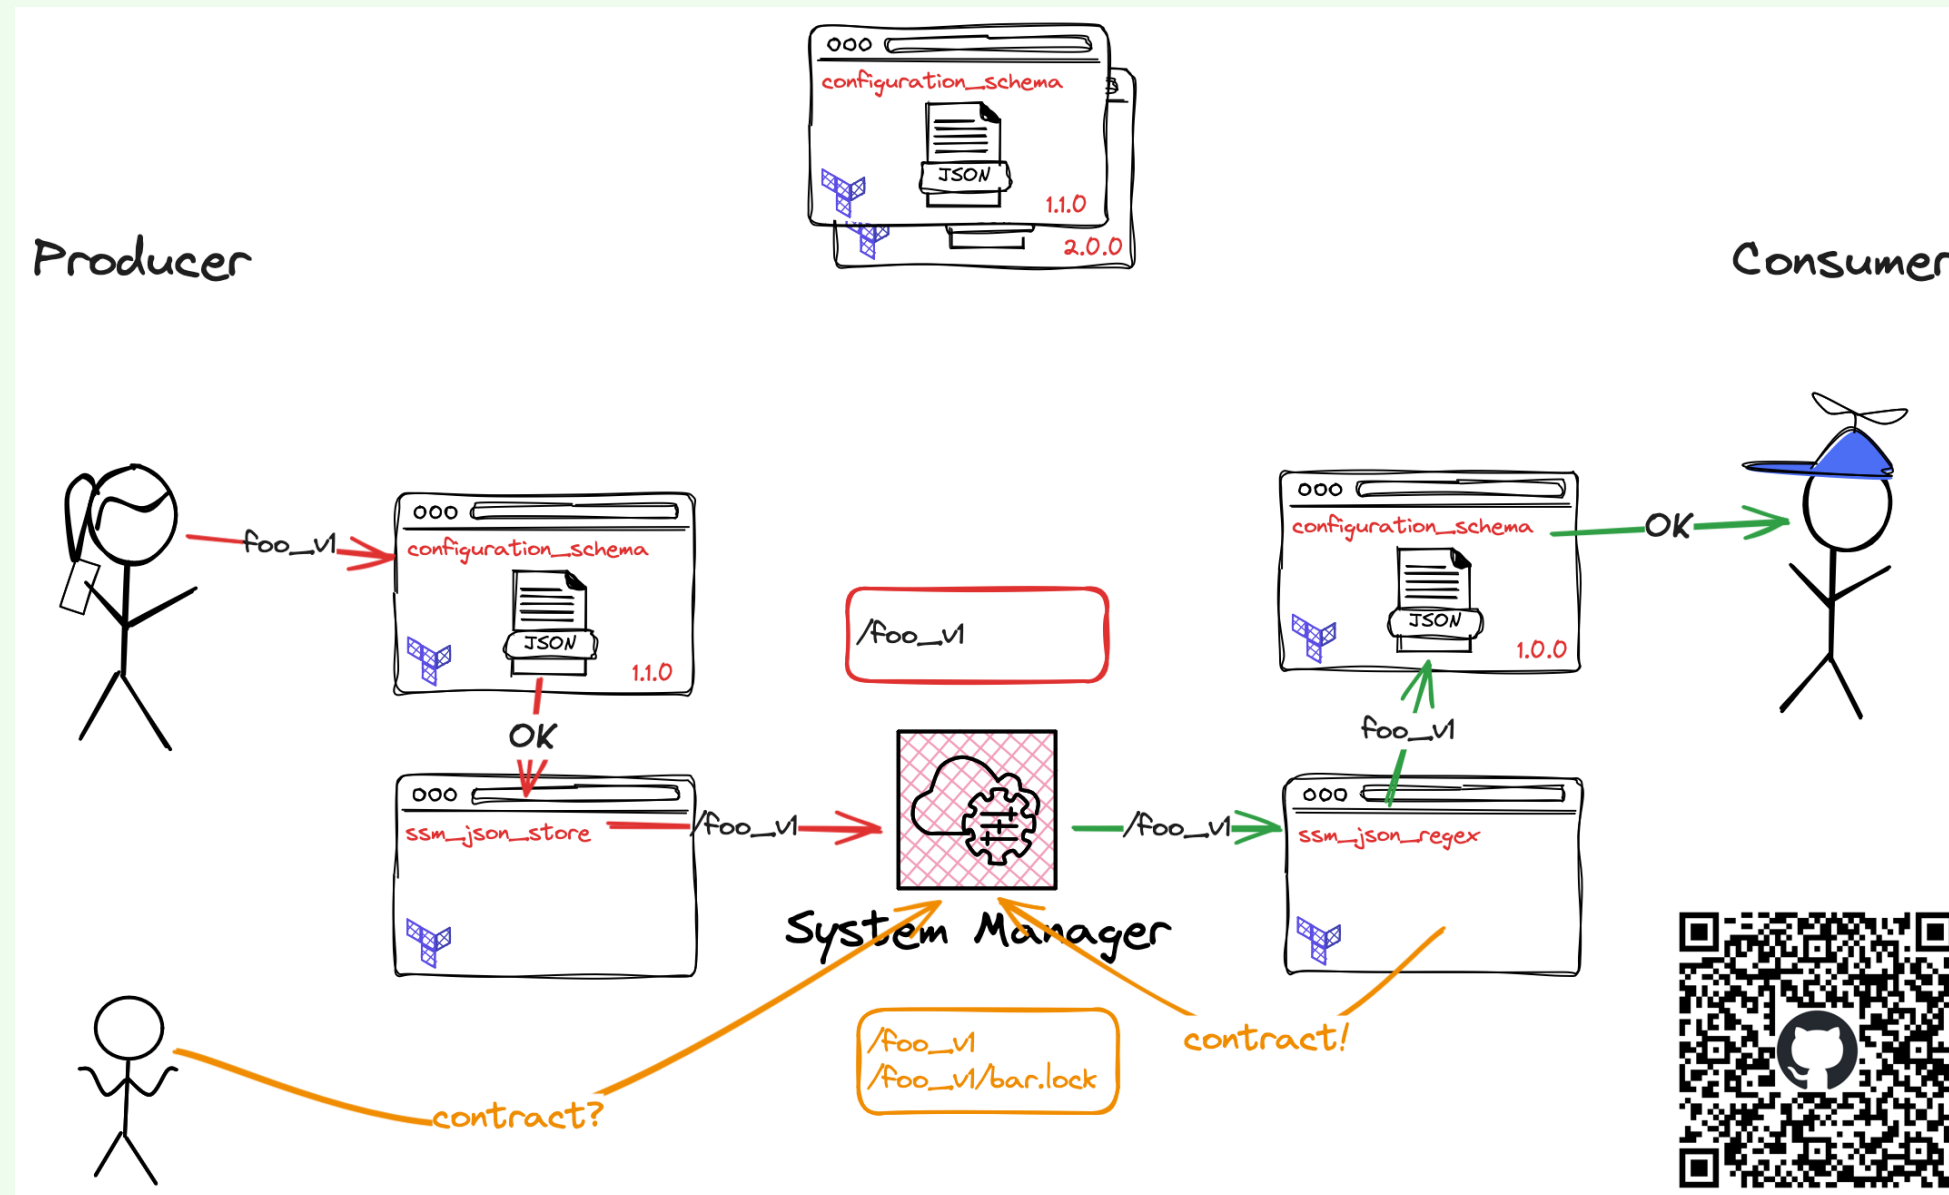
  sns_arn = nonsensitive(module.ssm_data.values["upstream"]["public"]["sns"]["foo"]["arn"])

  message_retention_seconds = 1209600
  redrive_policy = jsonencode({
    deadLetterTargetArn = module.dead_foo[0].arn
    maxReceiveCount = 5
  })
}
```

(1.3.5) Downsides of strong decoupling

- Data contracts between stacks
 - dependencies
 - versioning
- Dependencies of stacks
 - TF and Service code must be able to handle missing dependencies
 - reconciliation of TF stacks to check changed upstreams
 - eventually consistent
- Stack orchestration
 - state management should be centralised
 - stack execution should be in automation
- Permission management
 - for code changes (eg. CODEOWNERS)
 - for infrastructure changes
 - for accessing resources

(1.3.6) Soft data contract between stacks



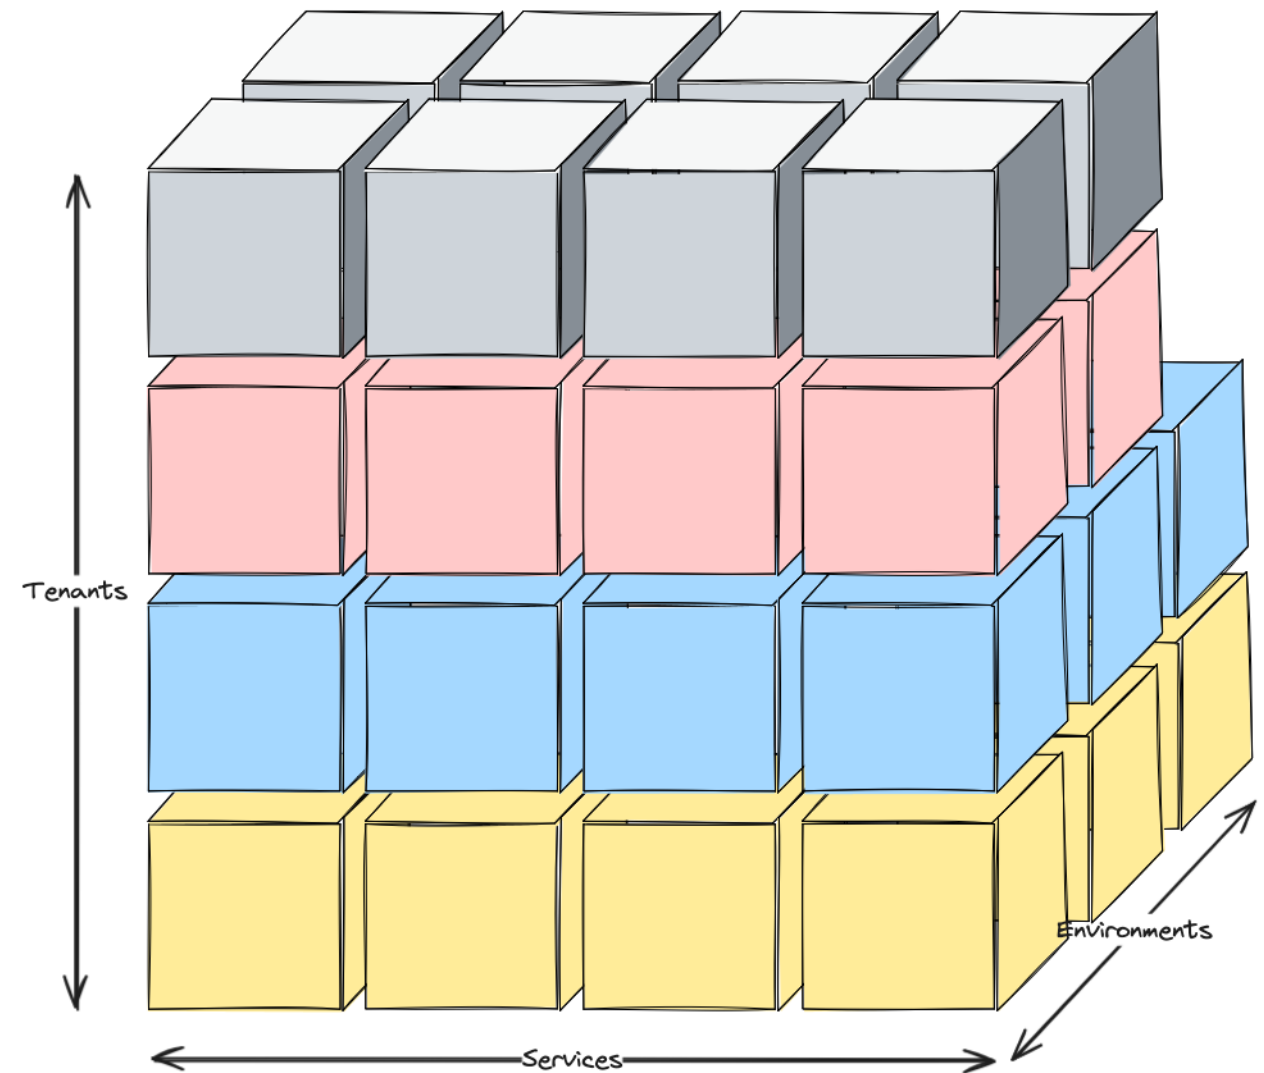
What's "reasonable scale", btw?

- we had 2 dimensions so far
 - number of TF stacks for x
 - number of environments for y
 - and a fixed number of tenants (1) for z
- let's expand the setup to multiple tenants
 - with this we'll get a real z dimension

total stacks = stacks * environments * tenants

To give some numbers: my client LYNQTECH runs ~100 microservices in at least 2 environments per tenant

for 5+ tenants - north of 1000 stacks 😊



(2) Terraform in GitOps

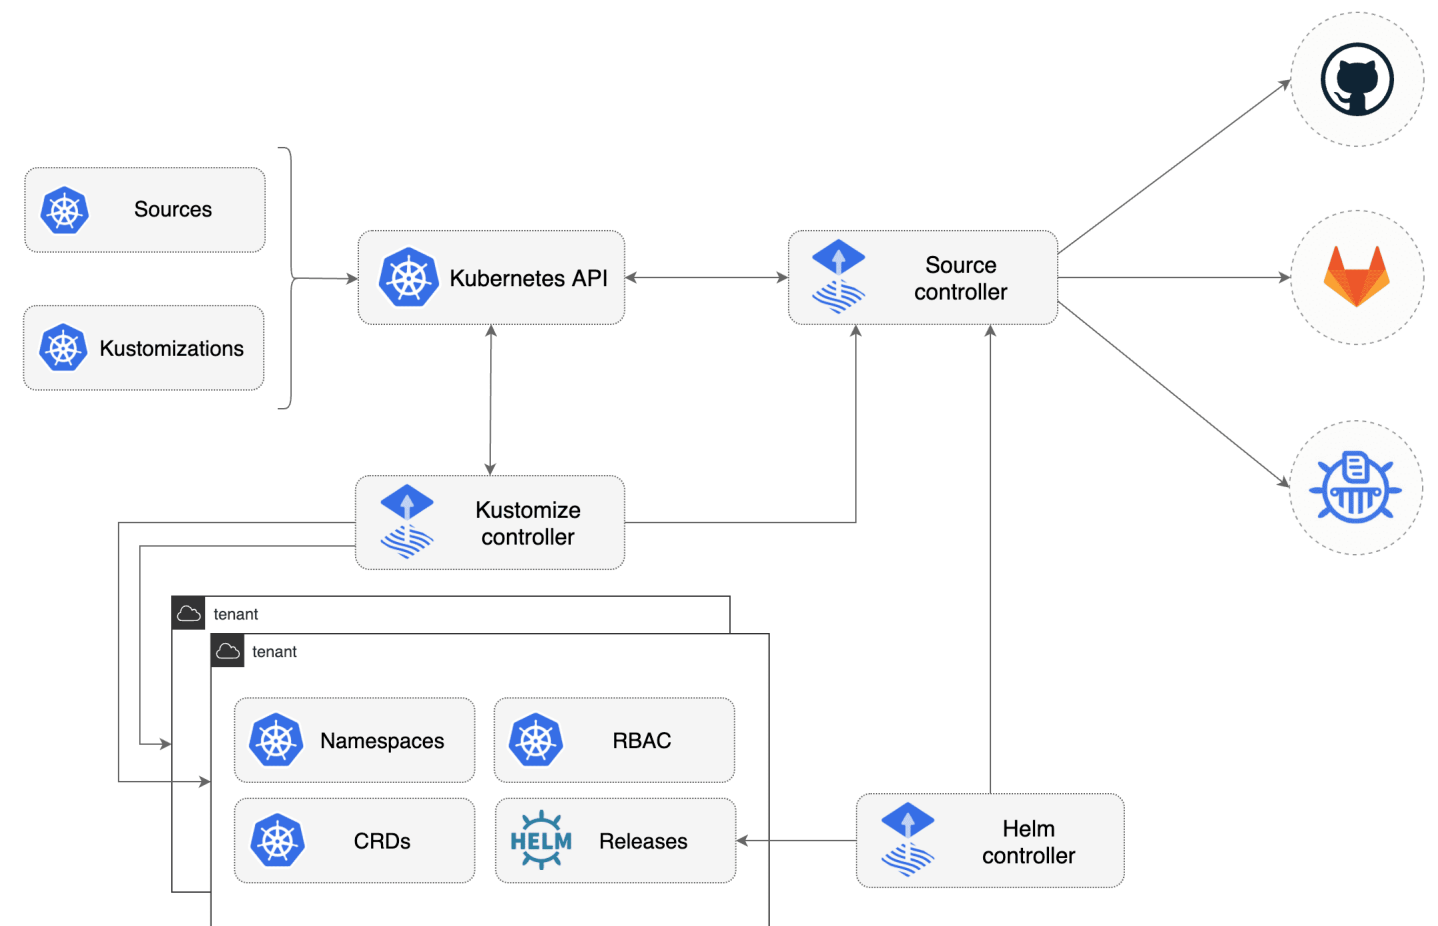
(2.1.1) FluxCD primer⁸

What is GitOps?

GitOps is an operational framework that takes DevOps best practices used for application development such as version control, collaboration, compliance, and CI/CD, and applies them to infrastructure automation.

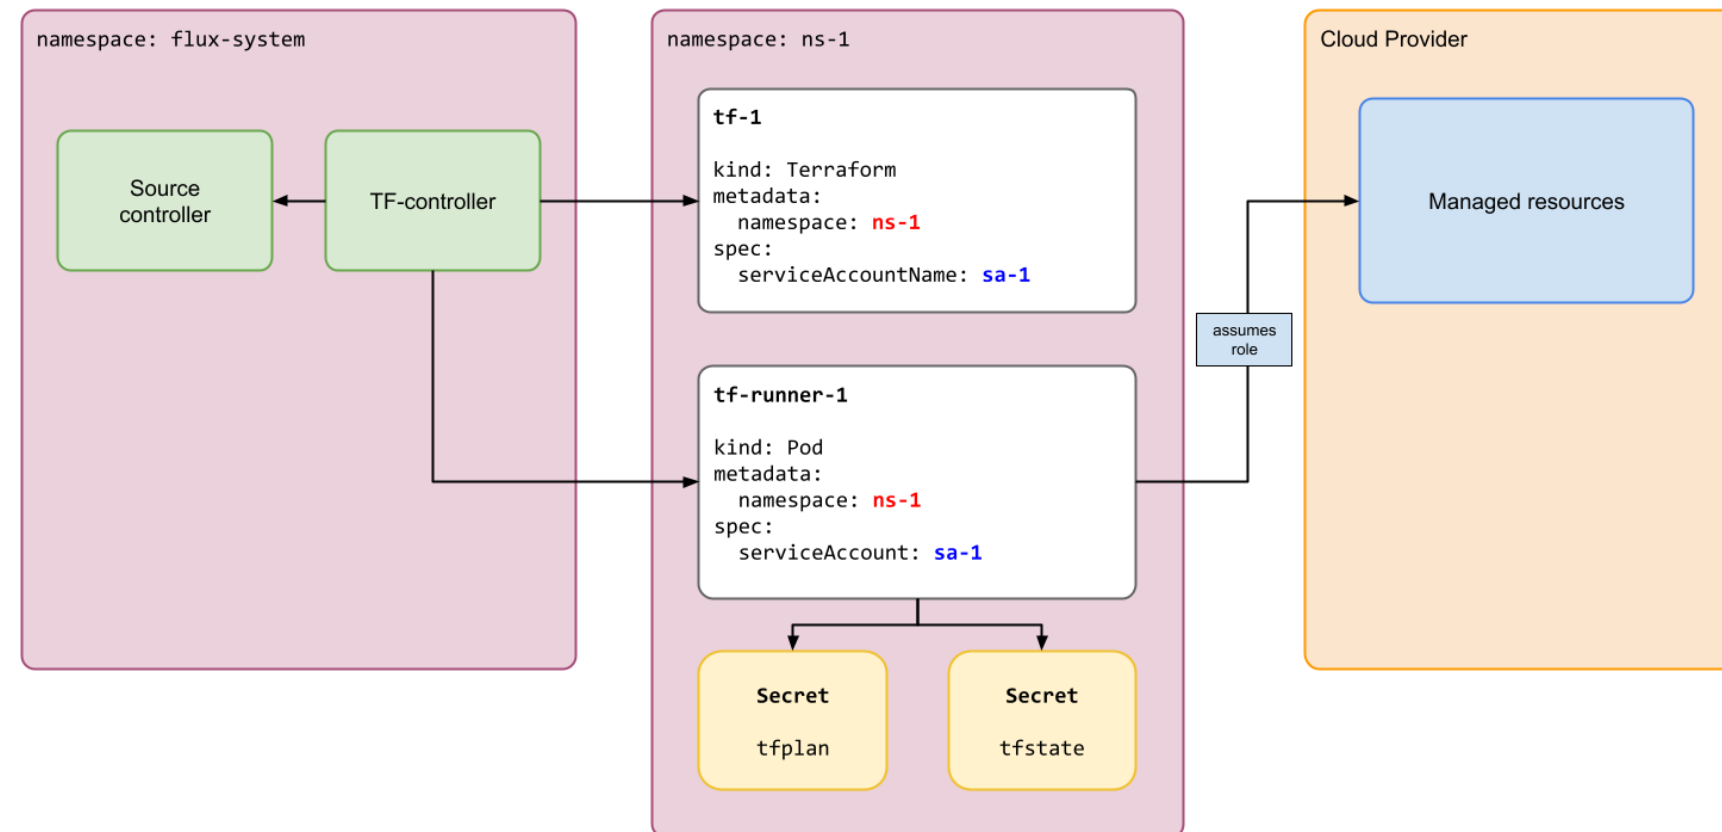
– <https://about.gitlab.com/topics/gitops/>

- In our context - **pull vs. push principle**
 - You don't care in *which environment* a stack runs in
 - They are ready for your stack and your code is pulled in (vs. pushed via a pipeline)



⁸ <https://fluxcd.io/flux/components/>

(2.1.2) Weaveworks TF Controller⁹ ¹⁸



⁹ <https://github.com/weaveworks/tf-controller>

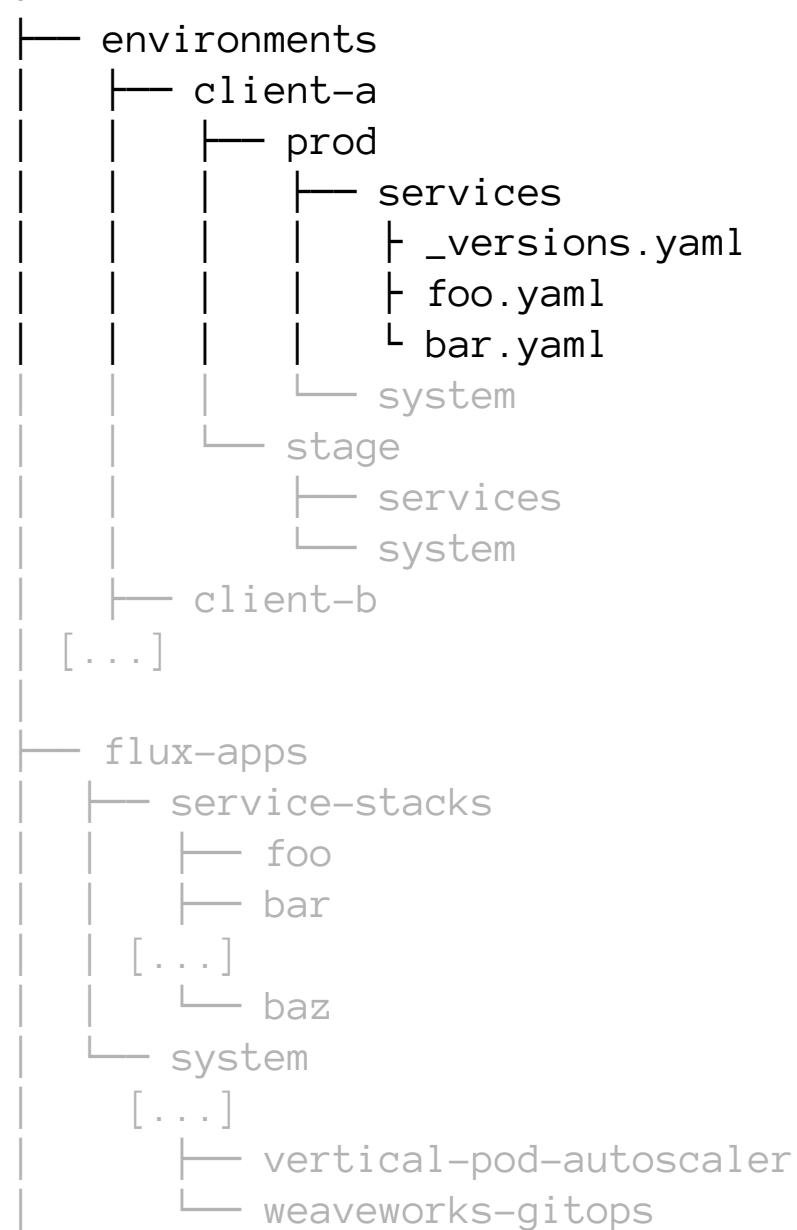
¹⁸ **Please note:** As the `tf-runner` ServiceAccount is usually very powerful, do not run it in an accessible namespace!

(2.2.1) Structure of central FluxCD configuration

- Each environment must be configurable individually
 - has its own entry point for FluxCD
 - this allows for configuration of deployed services
- For audit reasons, production environments must use fixed service versions, others can use semantic versioning
- Flux applications must be DRY
 - do not c&p code
 - implication: no individual configuration of apps
- in the central Flux repo there are **NO** variables, parameters etc. pp.
 - we only document **the intent** to run a service
 - self-configuration happens inside of an environment
- Use of OCI-based registries for sources only
 - everything as a final artefact
 - `flux push artifact`¹² is your friend

¹² https://fluxcd.io/flux/cmd/flux_push_artifact/

FluxCD as an App of Apps system



from the perspective of an individual FluxCD installation

- (0) - cloud and runtime is set up
 - provide data for stacks to become conscious
- (1) - load environment
 - primary Flux app
 - references all secondary service Flux apps
 - includes the version tracking ConfigMap
- (2) - load service Flux apps
 - contains relevant manifests
 - eg. OCI Sources, Terraform, Kustomization
- (3) - apply individual service apps

Primary Flux app

Applications > fluxcd-environment 🔔 🔍 👤

SYNC with Source without Source ▶ ⏸ ▶ Applied Revision: main@sha1:482... Last Updated: 4 minutes ago

✔ Applied revision: main@sha1:482acb06d5521f962959003db848ea56b533ff3

✔ All workloads are passing health checks

> More Information

DETAILS EVENTS **GRAPH** DEPENDENCIES YAML VIOLATIONS

```
graph TD; A["fluxcd-environment  
GitRepository  
flux-system"] --- B["✔ fluxcd-environment  
Kustomization  
flux-system"]; B --- C["✔ flux-system  
Kustomization  
flux-system"]; B --- D["✔ flux-system  
Kustomization  
flux-system"]; B --- E["✔ devops-sample-app  
Kustomization  
flux-system"]; B --- F["✔ flux-system  
Kustomization  
flux-system"];
```


Secondary Flux app

Applications > devops-sample-app

Applied Revision: main@sha1:482... Last Updated: 23 minutes ago

with Source without Source

Applied revision: main@sha1:482acb06d5521f962959003db848ea56b533ff3

All workloads are passing health checks

More Information

DETAILS EVENTS **GRAPH** DEPENDENCIES YAML VIOLATIONS

```
graph TD; A["fluxcd-environment  
GitRepository  
flux-system"] --> B["devops-sample-app  
Kustomization  
flux-system"]; B --> C["devops-sample-app  
Terraform  
flux-system"]; B --> D["devops-sample-app-iac  
OCIRepository  
flux-system"]; B --> E["devops-sample-app  
HelmRelease  
flux-system"];
```

85%

(2.3.1) Post build variable substitution¹⁰

- FluxCD's unique possibility to replace variables in rendered manifests before apply
- in FluxCD repo
 - environment specific `_versions.yaml` becomes `service-versions ConfigMap`
 - satisfies the "fixed versions" requirement
- In underlying IaC - basic environment information for a TF stack are written
 - base ConfigMap provides `client`, `environment` and other data
 - to form the path for Terraform state file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: service-versions
data:
  version_foo: "2.5.0"
  version_foo_tf: "~ 0.1.0-0"
  version_vertical_pod_autoscaler: "~> 9.0.0"
  version_vertical_pod_autoscaler_tf: "~ 0.1.0"
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: init
data:
  clientId: "tenant-a"
  domain: "stage.tenant-a.tld"
  environment: "stage"
  environmentClass: "non-prod"
  region: "eu-central-1"
```

¹⁰ <https://fluxcd.io/flux/components/kustomize/kustomizations/#post-build-variable-substitution/>

(2.3.2) usage example

```
apiVersion: source.toolkit.fluxcd.io/v1beta2
kind: OCIRepository
metadata:
  name: foo-iac
spec:
  interval: 5m
  provider: aws
  ref:
    semver: "${version_foo_tf}"
  url: oci://xxx.dkr.ecr.eu-central-1.amazonaws.com/iac/foo
```

```
apiVersion: infra.contrib.fluxcd.io/v1alpha2
kind: Terraform
metadata:
  name: foo
spec:
  backendConfig:
    customConfiguration: |
      backend "s3" {
        region          = "${region}"
        bucket           = "terraform-states"
        key              = "${clientId}/${environment}/stacks/foo.tfstate"
        role_arn         = "arn:aws:iam::xxx:role/tf-${clientId}-${environment}"
        dynamodb_table = "terraform-states-locks"
        encrypt          = true
      }
  sourceRef:
    kind: OCIRepository
    name: foo-iac
  vars: []
```

(2.4.1) Configuration Management

- (Terraform) code is agnostic of environments
- strict division of concerns between cloud and runtime environment
 - Helm/Kustomize - Runtime (Kubernetes)
 - Terraform - Cloud
- Each **Cloud** and **Runtime environment** allow a stack to become conscious
 - Cloud: SSM data base; Runtime: ConfigMap init
 - 🖱️ pull of configuration vs. push
- per **code stack** - data are baked into artifact
 - terraform - single configuration.tf
 - kustomize - separate overlay directories
 - helm - separate values.yaml

(2.4.2) Example

```
locals {
  service      = "foo"
  squad       = "bar"
  domain_name = module.ssm_data.values["base"]["domain"]
  cluster_name = module.ssm_data.values["base"]["cluster"]["name"]
  client       = nonsensitive(module.ssm_data.values["base"]["clientId"])
  environment  = nonsensitive(module.ssm_data.values["base"]["environment"])
  env_class    = nonsensitive(module.ssm_data.values["base"]["environment_class"])

  configuration = {
    default = {
      k8s_namespace      = local.service
      k8s_sa_name         = local.service
      rds_instance_class = "db.t4g.medium"
    }
    client_a = {
      stage = {}
    }
    environment_classes = {
      non-prod = {}
      prod = {
        rds_instance_class = "db.r6g.medium"
      }
    }
  }
}

# choose the right configuration based on
# client/environment/environment class or simply defaults
selected_configuration = merge(
  local.configuration["default"],
  try(local.configuration[local.client][local.environment], {})
)
}
```

```
# get the central SSM config parameters
module "ssm_data" {
  source = "registry.example.com/foo/ssm_full_json_store/aws"
  version = "0.3.1"

  path                = var.config_map_base_path
  include_filter_regex = "(base|foo|bar)"
}

module "database" {
  source = "registry.example.com/foo/RDS/aws"
  version = "3.5.0"

  identifier      = local.service
  squad           = local.squad
  rds_engine_version = local.selected_configuration["rds_engine_version"]
  rds_instance_class = local.selected_configuration["rds_instance_class"]
  client_id       = local.client
  environment     = local.environment
  vpc_id          = module.ssm_data.values["base"]["aws"]["vpc_id"]
  subnet_ids     = module.ssm_data.values["base"]["aws"]["subnet_public_ids"]
  # [...]
}
```

(2.4.3) Connecting Cloud and Runtime

- remember: division of concerns - cloud and runtime
- Terraform stack writes structured data as JSON
- Runtime pulls in data via External Secrets Operator¹⁶
- Reloader watches and upgrades Pods with their associated data

```
module "ssm_service_data" {
  source = "registry.example.com/foo/ssm_json_store/aws"
  version = "1.0.2"

  path = "/configuration"
  name = "foo"
  data = {
    installed = true
    private = {
      database = {
        database_name      = module.database.databases
        database_username = module.database.database_username
        endpoint            = module.database.endpoint
        reader_endpoint    = module.database.reader_endpoint
        port                = module.database.cluster_port
      }
    }
    public = {}
  }
}
```

```
apiVersion: external-secrets.io/v1beta1
kind: ExternalSecret
metadata:
  name: foo-secrets-ssm
spec:
  target:
    name: foo-secrets-ssm
  data:
    # [...]
    - remoteRef:
        key: /configuration/foo
        property: private.database.database_username
        secretKey: DATABASE_USER
    - remoteRef:
        key: /configuration/foo
        property: private.database.endpoint
        secretKey: DATABASE_HOST
---
kind: Deployment
metadata:
  annotations:
    reloader.stakater.com/auto: "true"
```

¹⁶ <https://external-secrets.io/>, [stakater/Reloader](https://stakater.com/)

(2.5) Specifics of TF-Controller

(2.5.1) Traffic

- each stack has its own tf-runner pod
 - **Decision:** no persistent pods between runs for security reasons (permissions of SA)
- Sizing example: terraform-provider-aws_5.31.0_darwin_arm64.zip = 84MB
- NAT costs (*AWS specific issue; GCP lowered egress costs to \$0 recently*)
 - reconcile every 30'
 - terraform init for each execution
 - 100 stacks * 48 runs/day * ~100MB providers * \$0,052/GB = **480GB/\$24,96 day/environment**
- **boring-registry** to the rescue 🎉
 - caching, pull-through proxy
 - Provider Network Mirror Protocol
- provider stored and delivered as S3 objects
- 😊 use S3 VPC endpoints

.terraformrc

```
credentials "my.terraform-registry.foo.bar" {  
  token = "7H151553CUr3!" # we are 1337  
}  
  
provider_installation {  
  network_mirror {  
    url = "https://my.terraform-registry.foo.bar/v1/mirror/"  
    include = ["*/*"]  
  }  
}
```

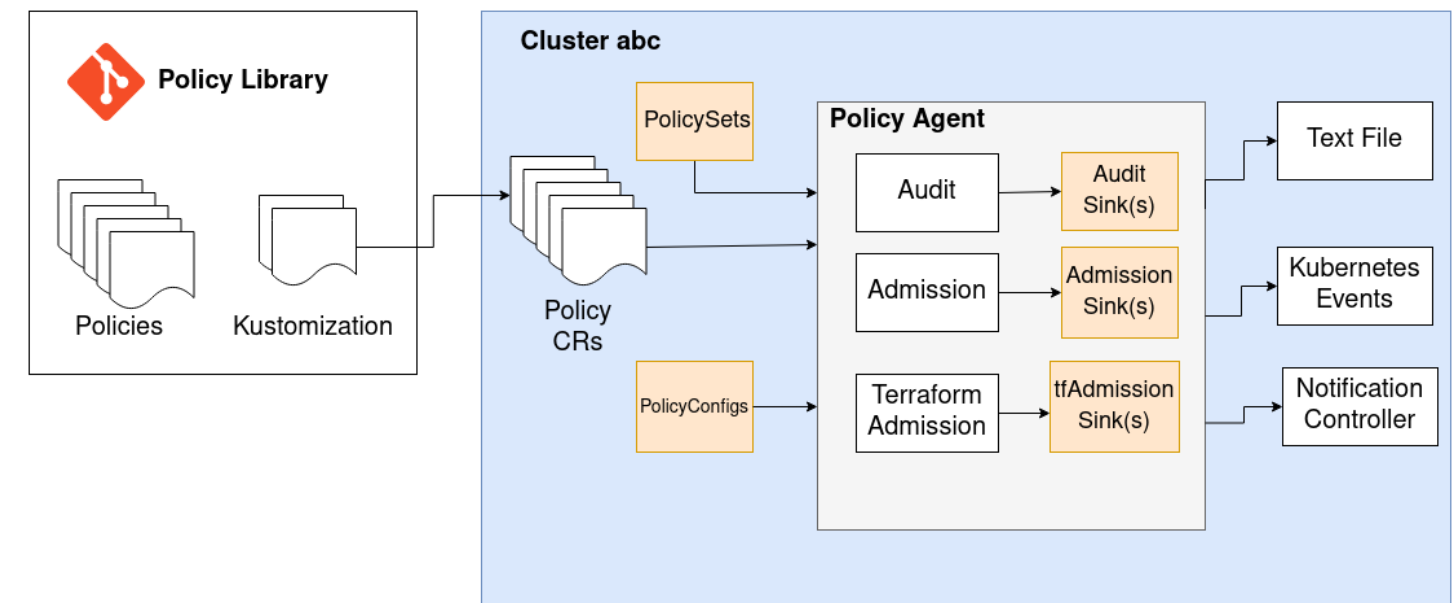
(2.5.2) Kubernetes resources

- each reconcile cycle triggers one tf-runner pod per stack
- each tf-runner pods consumes
 - ~800m CPU
 - ~150M Memory
- This would spawn a lot of machines at times
- using k8s limits based on priorityClass

```
apiVersion: scheduling.k8s.io/v1
description: used to limit the number of terraform runners
kind: PriorityClass
metadata:
  name: terraform
value: 0 # same priority as everybody else
---
apiVersion: v1
kind: ResourceQuota
metadata:
  name: terraform-runners
spec:
  hard:
    pods: "10"
  scopeSelector:
    matchExpressions:
    - operator: In
      scopeName: PriorityClass
    values:
    - terraform
```

(2.6) Weave Policy Engine¹⁷

- based on Rego and similar to Open Policy Agent
- **Goal:** auto approve Terraform changes
 - **Decision:** no destroy/recreate
 - **Decision:** no direct IAM resources (only via controlled modules)
- ⚠ **not an easy task - talk of its own**



¹⁷ [Weave Policy Engine](#), [Integrate TF Controller with Flux Receivers and Alerts](#), [Open Policy Agent](#)

(2.7) Weave GitOps UI¹¹

aka - the missing FluxCD UI

The screenshot displays the Weave GitOps UI interface, which is split into two main panels. The left panel, titled 'Applications', features a sidebar with navigation options: APPLICATIONS (selected), SOURCES, IMAGE AUTOMATION, FLUX RUNTIME, and DOCS. The main content area shows a table of applications with columns for NAME, KIND, NAMESPACE, TENANT, SOURCE, and STATUS. Two applications are listed: 'flux-system' (Kustomization) and 'ww-gitops' (HelmRelease), both in a 'Ready' status. The right panel, titled 'Flux Runtime', shows the cluster's Flux version (v0.34.0) and a table of controllers. The controllers table includes columns for NAME, STATUS, NAMESPACE, and IMAGE. Four controllers are listed: helm-controller, kustomize-controller, notification-controller, and source-controller, all in a 'Ready' status. The footer of the right panel provides contact information and version details: 'Need help? Contact us at sales@weave.works Kubernetes: v1.25.3 Flux: v0.34.0 Weave GitOps: v0.16.0 © 2023 Weaveworks'.

NAME	KIND	NAMESPACE	TENANT	SOURCE	STATUS
flux-system	Kustomization	flux-system	-	flux-system	Ready
ww-gitops	HelmRelease	flux-system	-	flux-system-ww-gitops	Ready

NAME	STATUS	NAMESPACE	IMAGE
helm-controller	Ready	flux-system	ghcr.io/fluxcd/helm-controller:v0.24.0
kustomize-controller	Ready	flux-system	ghcr.io/fluxcd/kustomize-controller:v0.28.0
notification-controller	Ready	flux-system	ghcr.io/fluxcd/notification-controller:v0.26.0
source-controller	Ready	flux-system	ghcr.io/fluxcd/source-controller:v0.29.0

¹¹ <https://github.com/weaveworks/weave-gitops> and <https://docs.gitops.weave.works/>

(3.0) Is it production ready?

- tf-controller is sometimes uncertain about the state
- slow development of tf-controller, ~~thank you HashiCorp~~
 - in principle ready for OpenTofu¹³
 - the talk uses features from a pre-release¹⁴
- observability is not ideal
 - eg. finding all Terraform Manifests, which have a pending plan

Be honest, where are you in the project?

- In the middle of cutting the large TF stacks
 - 🙌 very useful tool: [minamijoyo/tfmigrate](https://github.com/minamijoyo/tfmigrate)
- Automatic approvals are yet to come
- Branch Planner needs to be implemented to enable full developer ownership
- after IaC migration, services move to FluxCD as well

¹³ <https://www.opentofu.org/>

¹⁴ <https://github.com/weaveworks/tf-controller/releases/tag/v0.16.0-rc.3>

(3.1) Why not the BACK stack¹⁵?

- **Backstage (B)**: A self-service portal to empower developers
- **Argo CD (A)**: A GitOps-based continuous delivery (CD) tool for streamlined software delivery.
- **Crossplane (C)**: A universal control plane simplifying self-service infrastructure provisioning through abstractions.
- **Kyverno (K)**: A Policy as Code (PaC) tool
- existent Terraform stack and knowledge did not justify re-write of IaC
- Crossplane is bound to 1 kubernetes cluster (state in etcd) where Terraform is bound to a state file
- Introduction of Backstage was out of scope
- ArgoCD vs. FluxCD
 - ArgoCD's handling of Helm charts (templated and applied)
 - TF-Controller as part of FluxCD eco-system
 - ArgoCD has UI and concept of multi-cluster baked in
- Kyverno "runs as a dynamic admission controller" can not be used as a decision engine

¹⁵ Introducing the BACK Stack! - <https://www.youtube.com/watch?v=SMIR12uwMLs>

(3.2) Downsides

- development and local testing of TF code is hard
 - possibly via Branch Planner
 - *only for Github sources*
- Terraform module registry - so batteries included for developers?
yes, kind of, but
 - Terraform understanding needed
 - it is hard to grock the stack data exchange concept
 - we provide template repositories, use case documentation
- TF-Controller: (un)interruptable pods needed (for writing states)
- missing UI (for TF-Controller) and Monitoring APIs
- implicit data contracts between Terraform stacks

(3.2.1) - An uncertain future

The image shows a composite of three web browser screenshots. The top screenshot is a LinkedIn post from Alexis Richardson, CEO of Weaveworks, dated 2 days ago. The post text reads: "Hi everyone I am very sad to announce - officially - that Weaveworks will be closing its doors and shutting down commercial operations. Customers and partners will be working with a financial trustee whom we shall announce soon." The middle screenshot is from TechCrunch, dated April 25, 2024, with the headline "Cloud native container management platform Weaveworks shuts its doors" and a "Register Now" button. The bottom screenshot is from The Register, dated February 6, 2024, with the headline "GitOps pioneer Weaveworks unravels after funding fabric frays" and a sub-headline "Company burned through \$61.6M in investment".

https://www.linkedin.com/feed/update/urn:li:activity:7160295096825860096/

in Search Home My Network Jobs Messaging Notifications

Meet us at KubeCon Booth S55 Flux Booth

Alexis Richardson
CEO Weaveworks
View full profile

Alexis Richardson · Following
CEO Weaveworks
2d · Edited · 🌐

Hi everyone

I am very sad to announce - officially - that Weaveworks will be closing its doors and shutting down commercial operations. Customers and partners will be working with a financial trustee whom we shall announce soon.

https://techcrunch.com/2024/02/05/cloud-native-container-management-platform-weaveworks-shuts-its-doors/

Startups

Cloud native container management platform Weaveworks shuts its doors

TechCrunch Early Stage April 25, 2024 Boston, MA
Register Now

https://www.theregister.com/2024/02/06/weaveworks_folds/

SIGN IN / UP The Register

CXO 3

GitOps pioneer Weaveworks unravels after funding fabric frays

Company burned through \$61.6M in investment

2024-02-06 - <https://www.linkedin.com/feed/update/urn:li:activity:7160295096825860096/>

(3.3) Upsides

- all domains (code, kubernetes and cloud environment) follow the same pattern
 - same CI and CD
 - same artefact type (OCI)
 - similar release cycles
 - single entry point for Product Owners
- IaC runner can be replaced
 - TF-Controller is just **a** controlled terraform executor
 - migration to eg. Spacelift.io or others possible
 - break-the-glas scenario supported (manual stack execution)
- Terraform/OpenTofu eco-system can be reused
 - providers
 - knowledge and modules

(3.4) Thanksides

- LYNQTECH GmbH for granting permission to share information and code
 - 😊 LYNQTECH is hiring
<https://www.lynq.tech/jobs/>
- All colleagues who were and are part of this journey
- The FluxCD Community and WeaveWorks for their software



LinkedIn



Website



GitHub

(4.0) Architectural decisions

General FluxCD

- Each tenant environment must be configurable individually
 - For audit reasons, production envs must use fixed service versions
- Applications, in the central repo, must be DRY. No individual stacks.
- Use of OCI-based registries for sources only (exception: external Helm)
- Code is agnostic of environments and is not parameterised
 - Each **cloud (AWS)** and **runtime (Kubernetes) environment** allows a stack to become concious
 - kustomize style data baked into artifact
- Secrets synchronised via External-Secret Operator
- Kubernetes cluster should be treated as cattle

TF-Controller

- No vendor lock-in; re-usability of eco-system strong plus
 - Terraform providers
 - Terraform OSS modules
- No persistent pods between runs
- Aim for Auto approval for Terraform changes
 - no destroy/recreate
 - no direct IAM resources (only via controlled modules)
 - only approved top-level module sources

Image sources

1. FluxCD documentation - <https://fluxcd.io/flux/components/>
2. Weave GitOps // Terraform Controller documentation - <https://weaveworks.github.io/tf-controller/>
3. Weave GitOps // The Policy Ecosystem - <https://docs.gitops.weave.works/docs/policy/getting-started/>