

# **LESSONS LEARNED FROM RUNNING TERRAFORM AT REASONABLE SCALE**



Daniel Ciaglia

# UTILIZING FLUXCD, WEAVEWORKS TF-CONTROLLER AND BORING-REGISTRY

*Why easy, when we can make it complicated? – the unknown platform engineer*

~~~~

## USAGE/NAVIGATION

- n, p - next, previous slide
- o - overview
- f - fullscreen
- b - black out the presentation
- s - speaker view
-   - navigate on the top-level

# DANIEL CIAGLIA // *CONSULTANT*

- Freelance  
*since 2022*
- TIER Mobility SE  
*Director of Engineering*
- kreuzwerker GmbH  
*Principal Consultant*
- Bundesdruckerei GmbH  
*Senior Support Manager*
- *[some more]*
- SCUBA dive instructor
- AWS User Group Berlin  
co-organiser



LinkedIn



Website

# TODAY'S MENU

1. A typical Terraform stack evolution
2. Running Terraform in GitOps
3. Thoughts on the stack
4. Architectural Decision Records summary

# (1.1) TYPICAL TERRAFORM STACK EVOLUTION<sup>1</sup>

*Stack: Terraform root module<sup>2</sup>, tracked with 1 state file*

*Related: Highly recommend talk “Terraform: from zero to madness” by [@Timur Bublik](#)*

## (1.1.1) IN THE BEGINNING

- you start your project
- put everything in 1 directory
- maybe split files by broader domains.

```
.  
├── databases.tf  
├── vpc.tf  
├── main.tf  
├── outputs.tf  
└── terraform.tf
```

## (1.1.2) THE STAGING/PRODUCTION SPLIT

- oh well, you need a staging environment
- both environments are very much the same
- you refactor the code to be parameterised by variables
- you provide 2 `.tfvars` files

```
1 .
2 |— production.tfvars
3 |— staging.tfvars
4 |— databases.tf
5 |— vpc.tf
6 |— variables.tf
7 |— main.tf
8 |— terraform.tf
```

## (1.1.3) CODE REPETITION - I NEED MODULES

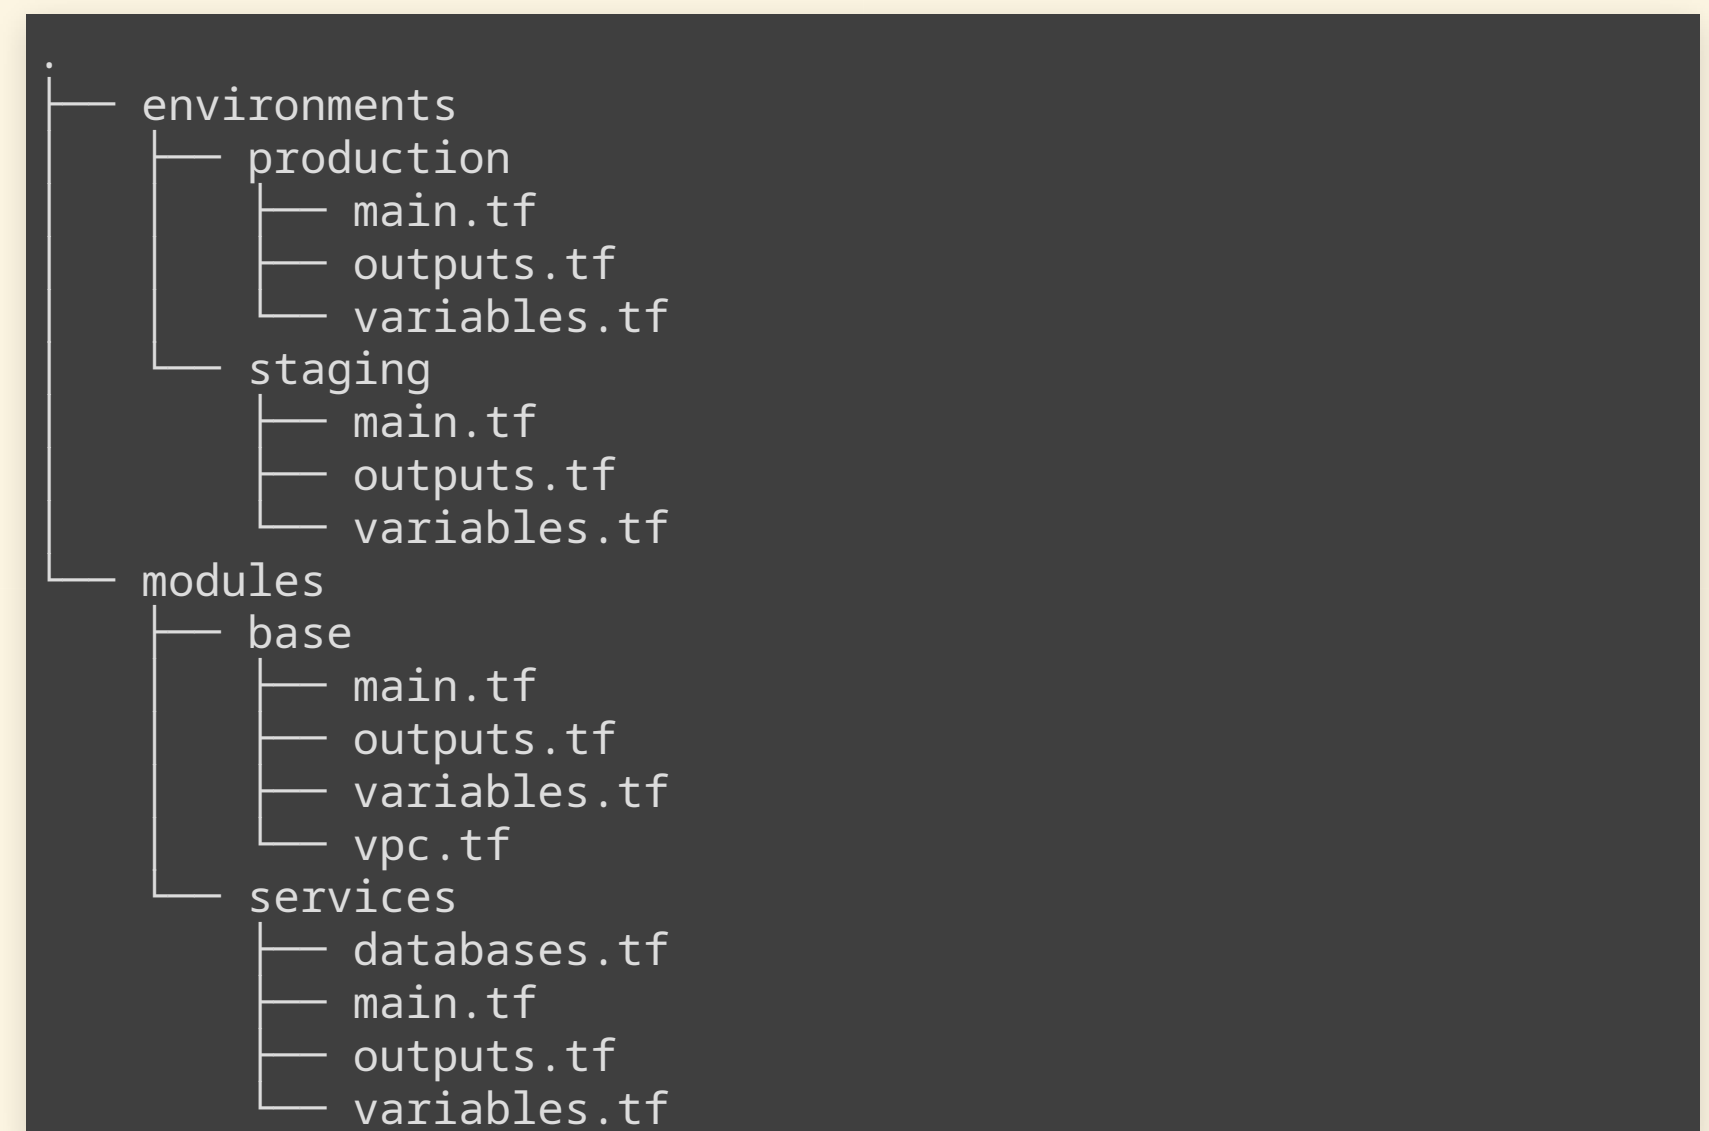
- you add more services and they need infra
- the infra is similar
- you want to keep the code DRY<sup>3</sup>
- you create a repo, codify best practices, tag them for versioning
- you pull in modules via git

```
# select a specific tag
module "rds" {
  source = "github.com/example/rds?ref=v1.2.0"
}
```



## (1.1.4) THE GREAT SEPARATION

- as the stack grows, the environments differ
- you start separating the code in larger blocks
  - the base environment
  - the services
- the code is pulled in as modules
- the `services` module receives output of base as input eg. `vpc_id` or `subnets`
- `terraform apply plan` is run manually still



# (1.1.5) FAST FORWARD

👉 At this point in time I joined the project 👈

## The situation

- as the stack grows further, the amount of resources does as well
- each run of `terraform plan -out plan` takes more and more time
- to review and apply changes for developers becomes a dayfilling job
- you start cheating by targeted apply
- you notice that the amount of files downloaded for each `terraform` step is enormous<sup>4</sup>
- you notice that git tags can not be used for semantic versioning (`version`)

## Possible solutions

- to address the versioning and data transfer issues - use a private Terraform module registry
- to address the runtime and ownership issue - split the stacks and let the teams handle them (DevOps style)

## (1.2) THE BORING-REGISTRY

- TIER Mobility developed their own “boring” Terraform registry without moving parts (hence the name)
  - Details to be found here: <https://github.com/boring-registry/boring-registry/>
  - The important feature for now is support for the [Module Registry Protocol](#)
- **You provide** a S3 bucket, module code and package it in CD via `./boring-registry upload --type s3 (some more flags) ./your-module`
- **You’ll get** semantic versioning

```
module "rds" {  
  source = "registry.example.com/acme/rds/aws"  
  version = "~> 0.1"  
}
```

# (1.3) SEPARATING THE SERVICE STACKS

## SOME ARCHITECTURAL DECISIONS

### Don't

- separate services along team borders<sup>5</sup> → teams and responsibilities change, always
- share states between services → there are secrets in there!<sup>6</sup> → read the docs of the [terraform\\_remote\\_state](#) data source!

### Do

- Layer your stacks - account, network, clusters and services
- 1 Terraform stack per service
  - good for least privilege access
  - place the Terraform code into the service repo
- run the TF stacks in automation
- use an indirect way to share information between stacks<sup>7</sup>

## (1.3.1) INDIRECT INFORMATION EXCHANGE

- use structured data → ideally JSON for `jsondecode()` and `jsonencode()`
- use whatever storage you prefer → SSM Parameter Store or S3

Code for 3 Terraform modules will be provided

- `s3_json_store` CRUD JSON data on S3
- `ssm_json_store` CRUD JSON data on SSM Parameter store
- `ssm_json_regex` read SSM parameter with regex

## (1.3.2) WRITE DATA (BASE SYSTEM)

```
1 module "ssm_service_data" {
2   source = "registry.example.com/foo/ssm_json_store/aws"
3   version = "~> 1.0.2"
4
5   path = "/configuration"
6   name = "base"
7   data = {
8     domain          = local.domain_name
9     environment     = local.environment
10    environmentClass = local.environmentClass
11    backup_plan     = local.backup_plan
12    networking = {
13      vpc_id          = module.base.vpc_default_id
14      subnet_database_ids = module.base.subnet_private_database_ids
15      subnet_k8s_ids   = module.base.subnet_private_k8s_ids
16    }
17    cluster = {
18      name          = module.eks.cluster_name
19      oidc_issuer_url = module.eks.cluster_oidc_issuer_url
20      oidc_provider_arn = module.eks.cluster_oidc_provider_arn
21    }
22  }
23 }
```

## (1.3.3) WRITE DATA (UPSTREAM)

```
1 module "ssm_service_data" {
2   source = "registry.example.com/foo/ssm_json_store/aws"
3   version = "~> 1.0.2"
4
5   path = "/configuration"
6   name = "upstream"
7   data = {
8     installed = true
9
10    private = {}
11    public = {
12      sns = {
13        "foo" = {
14          "arn" = module.sns_foo.arn
15          "name" = module.sns_foo.name
16        }
17        sqs = {
18          "bar" = {
19            "arn" = module.bar_queue.arn
20            "name" = module.bar_queue.name
21          }
22        }
23      }
24    }
25  }
26 }
```

## (1.3.4) READ DATA (DOWNSTREAM)

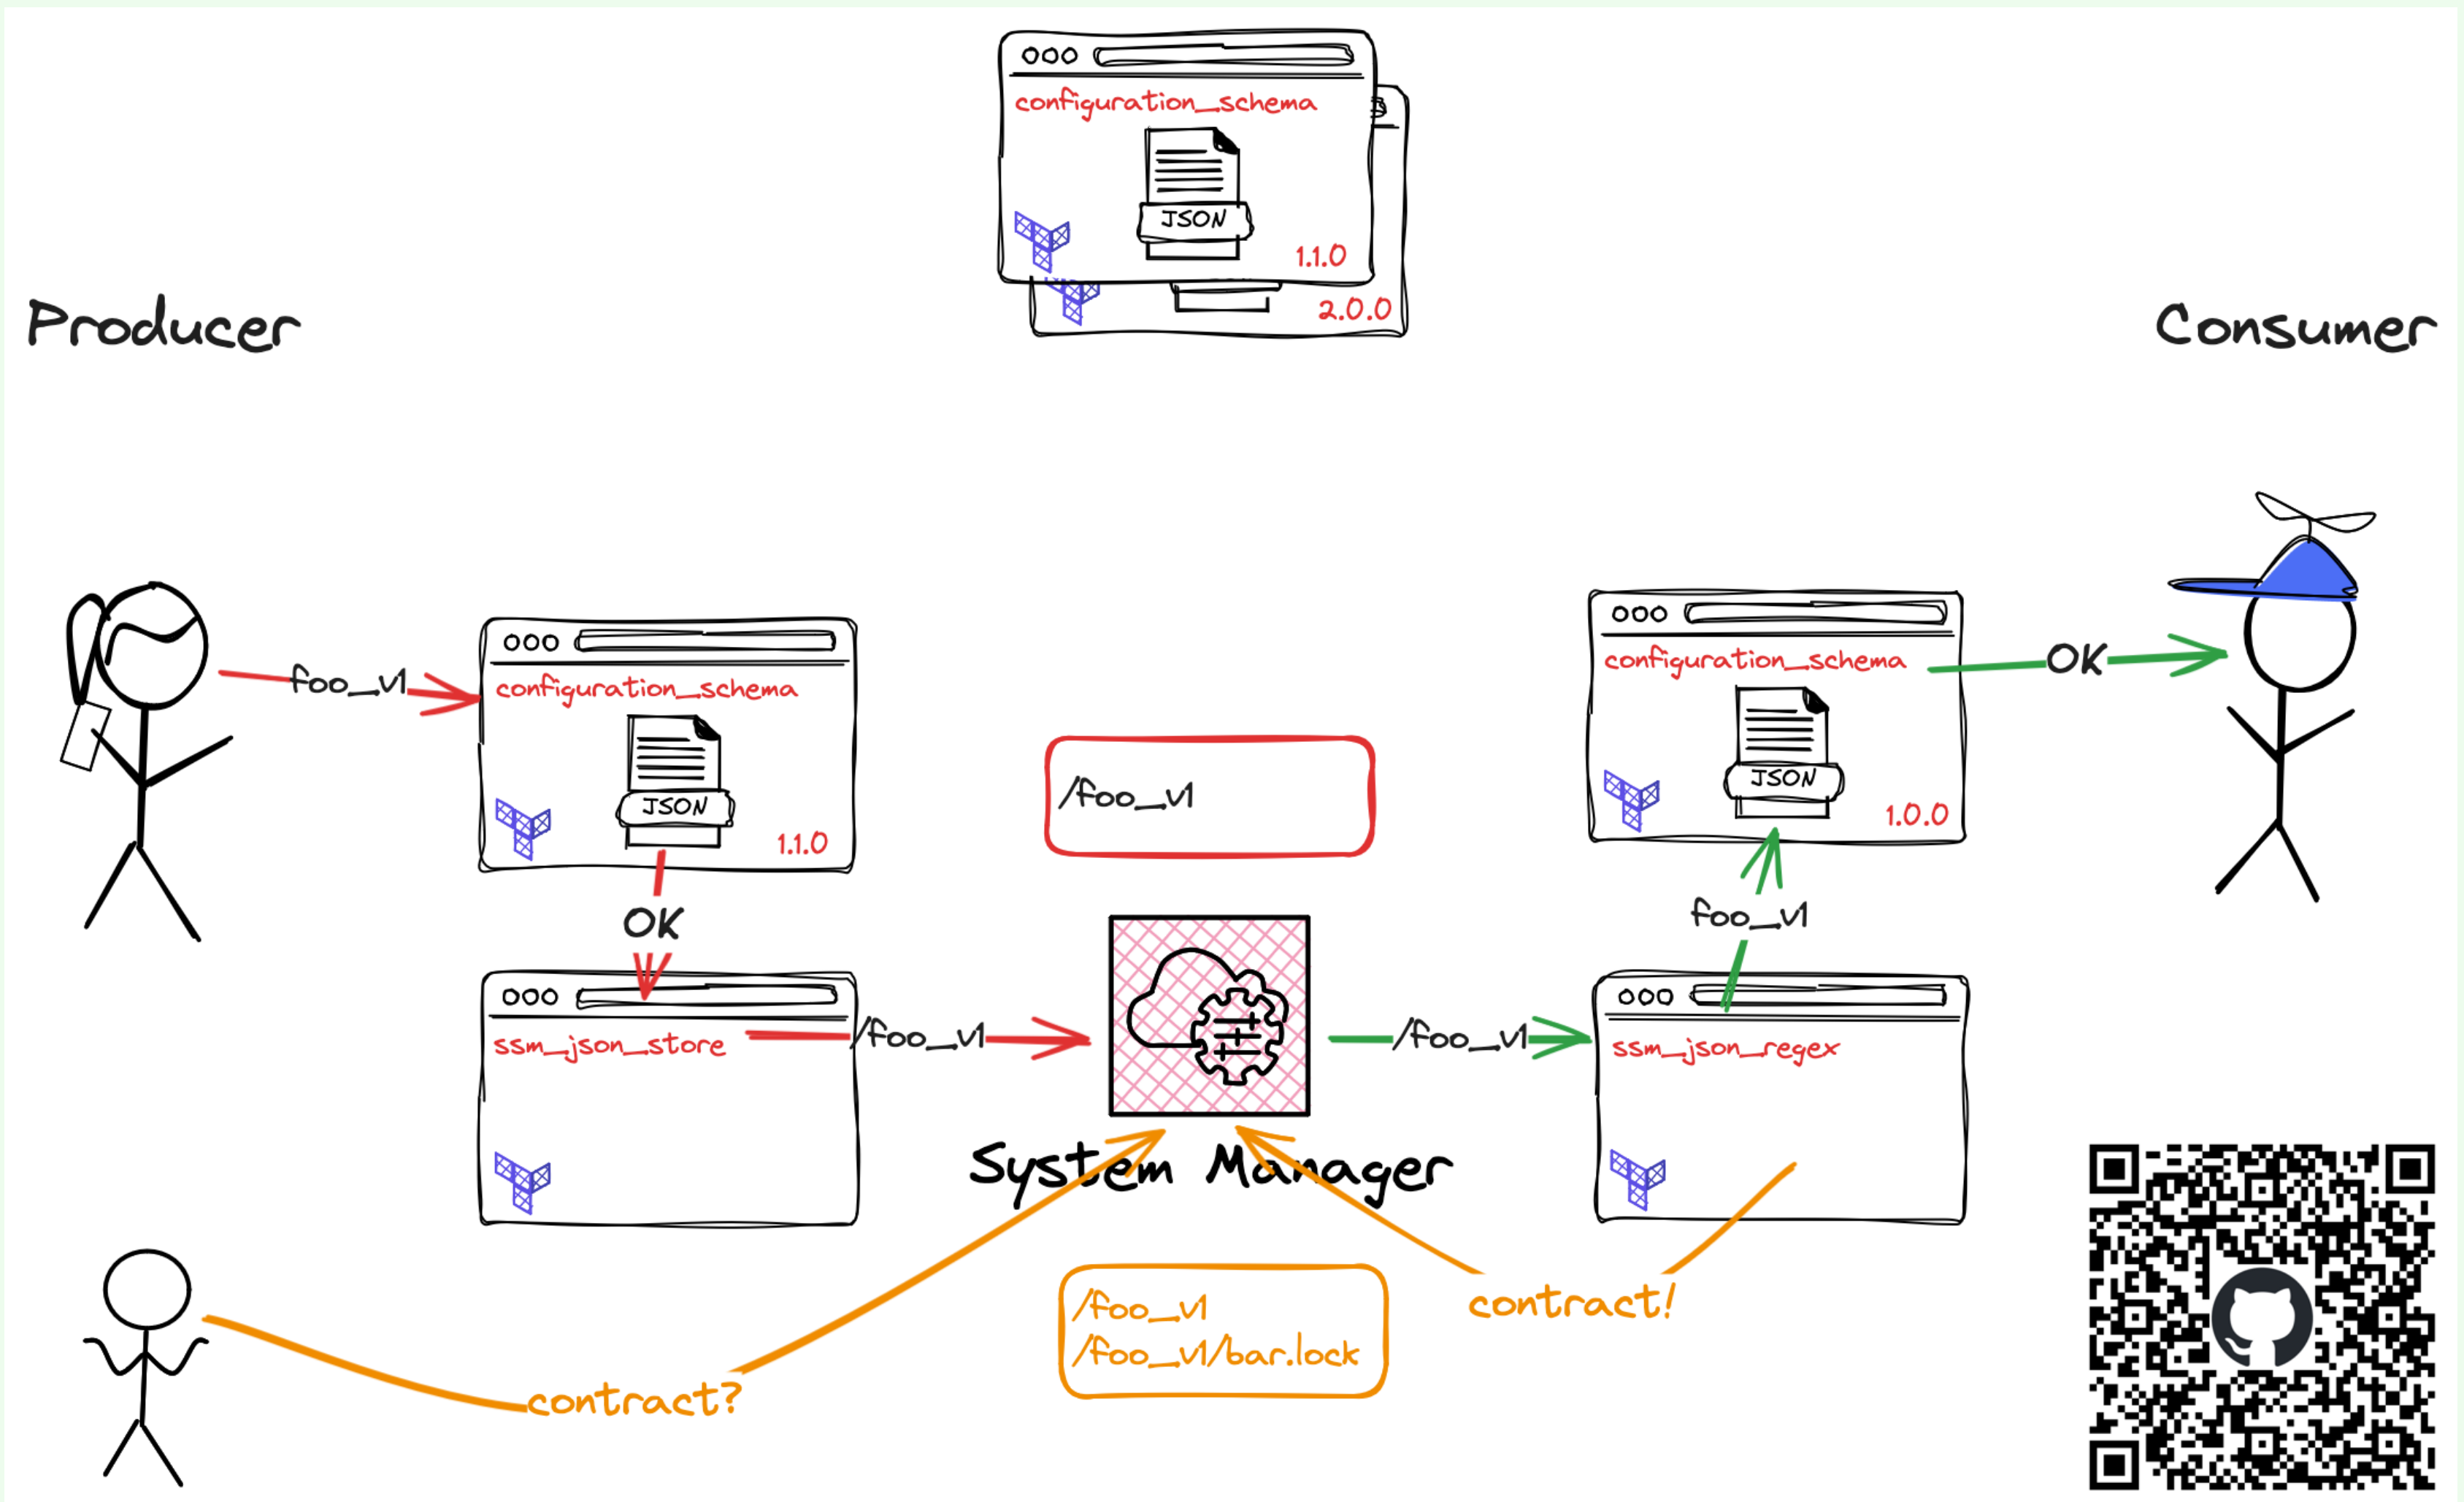
```
1 module "ssm_data" {
2   source = "registry.example.com/foo/ssm_json_store/aws"
3   version = "~> 0.1.0"
4
5   path = "/configuration"
6   include_filter_regex = "(base|upstream)"
7 }
8
9 module "sns_sqs_subscription_foo" {
10  count = try(module.ssm_data.values["upstream"]["installed"], false) ? 1 : 0
11  source = "registry.example.com/foo/sns_sqs_subscription/aws"
12  version = "~> 0.1"
13
14  sns_arn = nonsensitive(module.ssm_data.values["upstream"]["public"]["sns"]["foo"]["arn"])
15
16  message_retention_seconds = 1209600
17  redrive_policy = jsonencode({
18    deadLetterTargetArn = module.dead_foo[0].arn
19    maxReceiveCount = 5
20  })
21 }
```



## (1.3.5) DOWNSIDES OF STRONG DECOUPLING

- Data contracts between stacks
  - dependencies
  - versioning
- Dependencies of stacks
  - TF and Service code must be able to handle missing dependencies
  - reconciliation of TF stacks to check changed upstreams
  - eventually consistent
- Stack orchestration
  - state management should be centralised
  - stack execution should be in automation
- Permission management
  - for code changes (eg. CODEOWNERS)
  - for infrastructure changes
  - for accessing resources

# (1.3.6) SOFT DATA CONTRACT BETWEEN STACKS




# WHAT'S "REASONABLE SCALE", BTW?

- we had 2 dimensions so far
  - number of TF stacks for x
  - number of environments for y
  - and a fixed number of tenants (1) for z
- let's expand the setup to multiple tenants
  - with this we'll get a real z dimension

**total stacks = stacks \* environments \* tenants**

To give some numbers: my client [LYNQTECH](#) runs ~100 microservices in at least 2 environments per tenant for 5+ tenants - north of 1000 stacks 😊



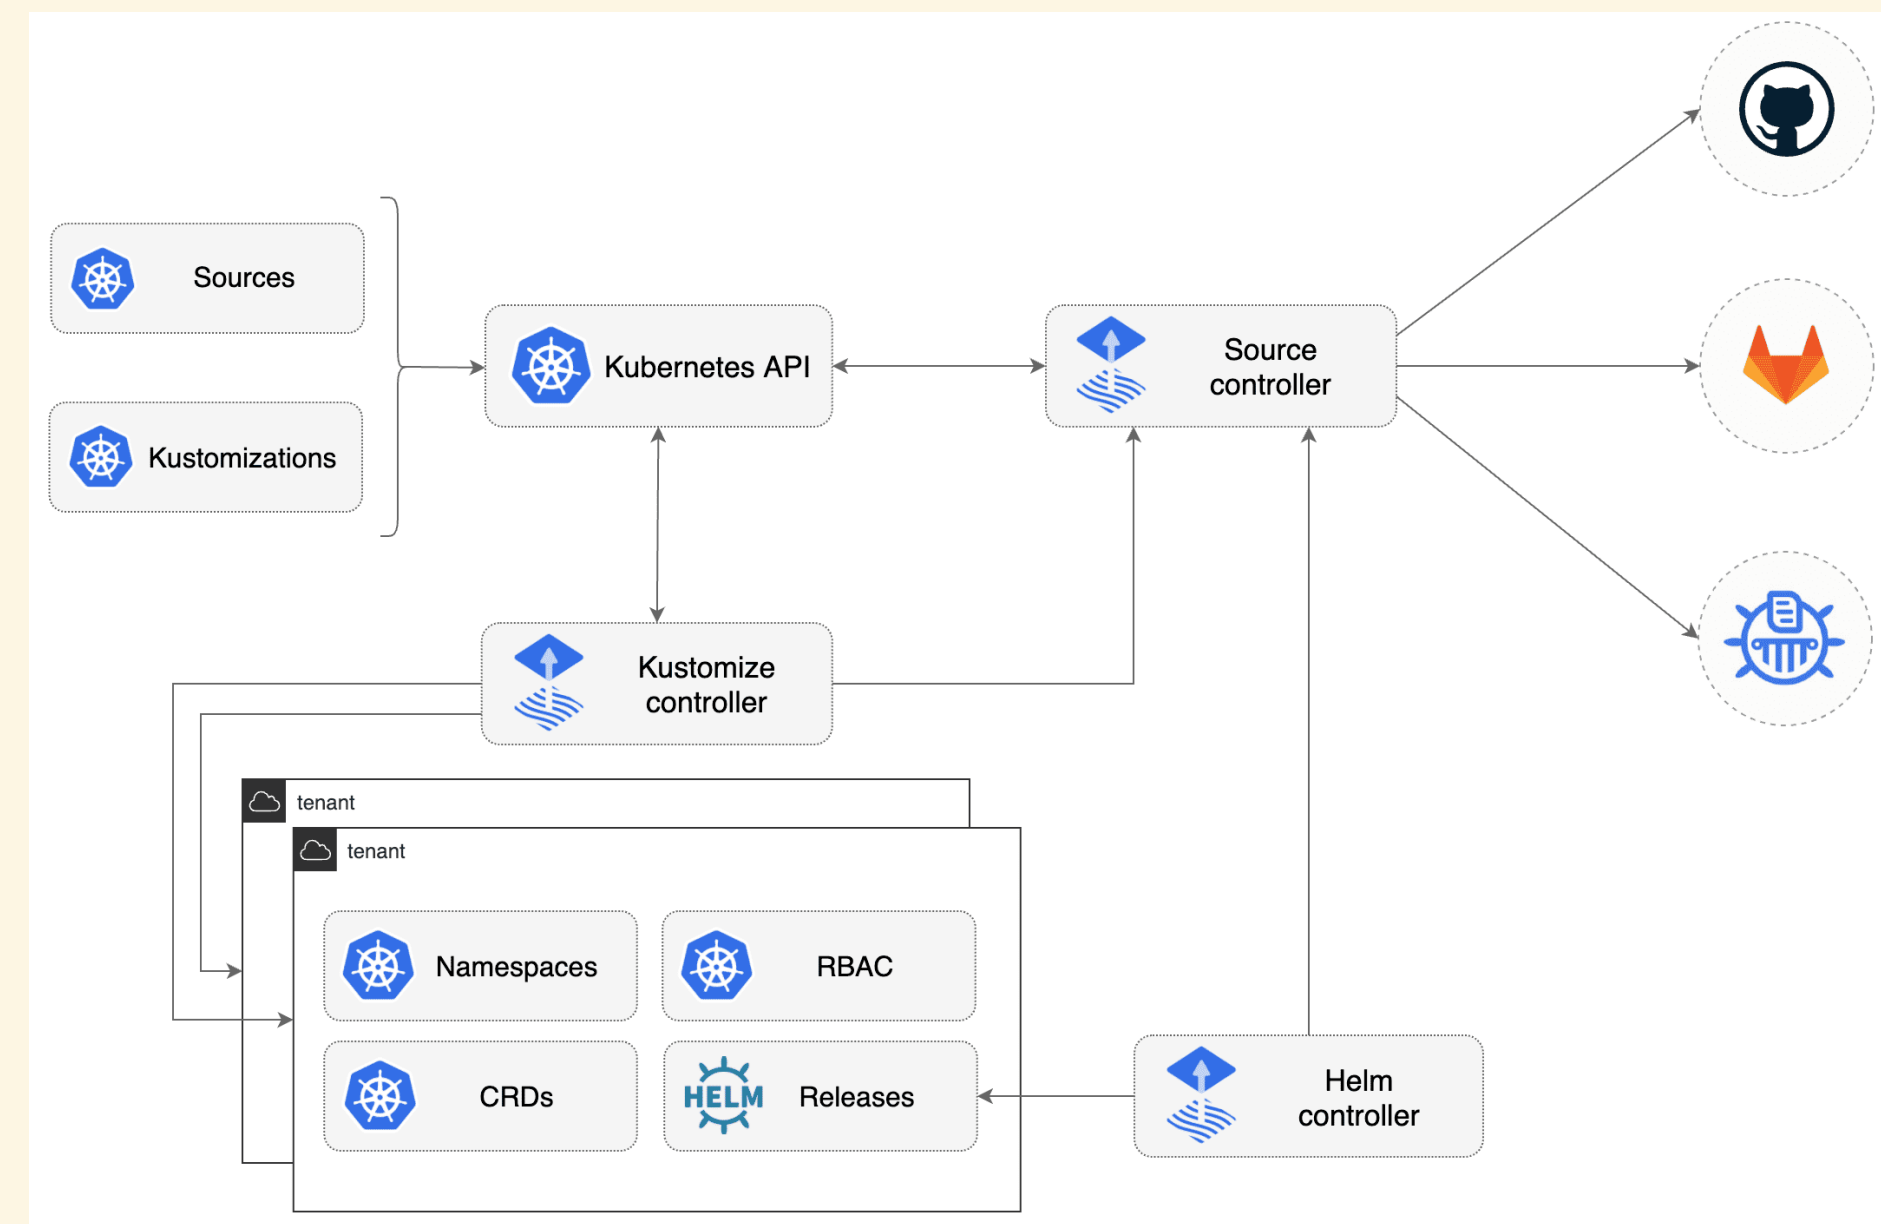
# **(2) TERRAFORM IN GITOPS**

## (2.1.1) FLUXCD PRIMER<sup>8</sup>

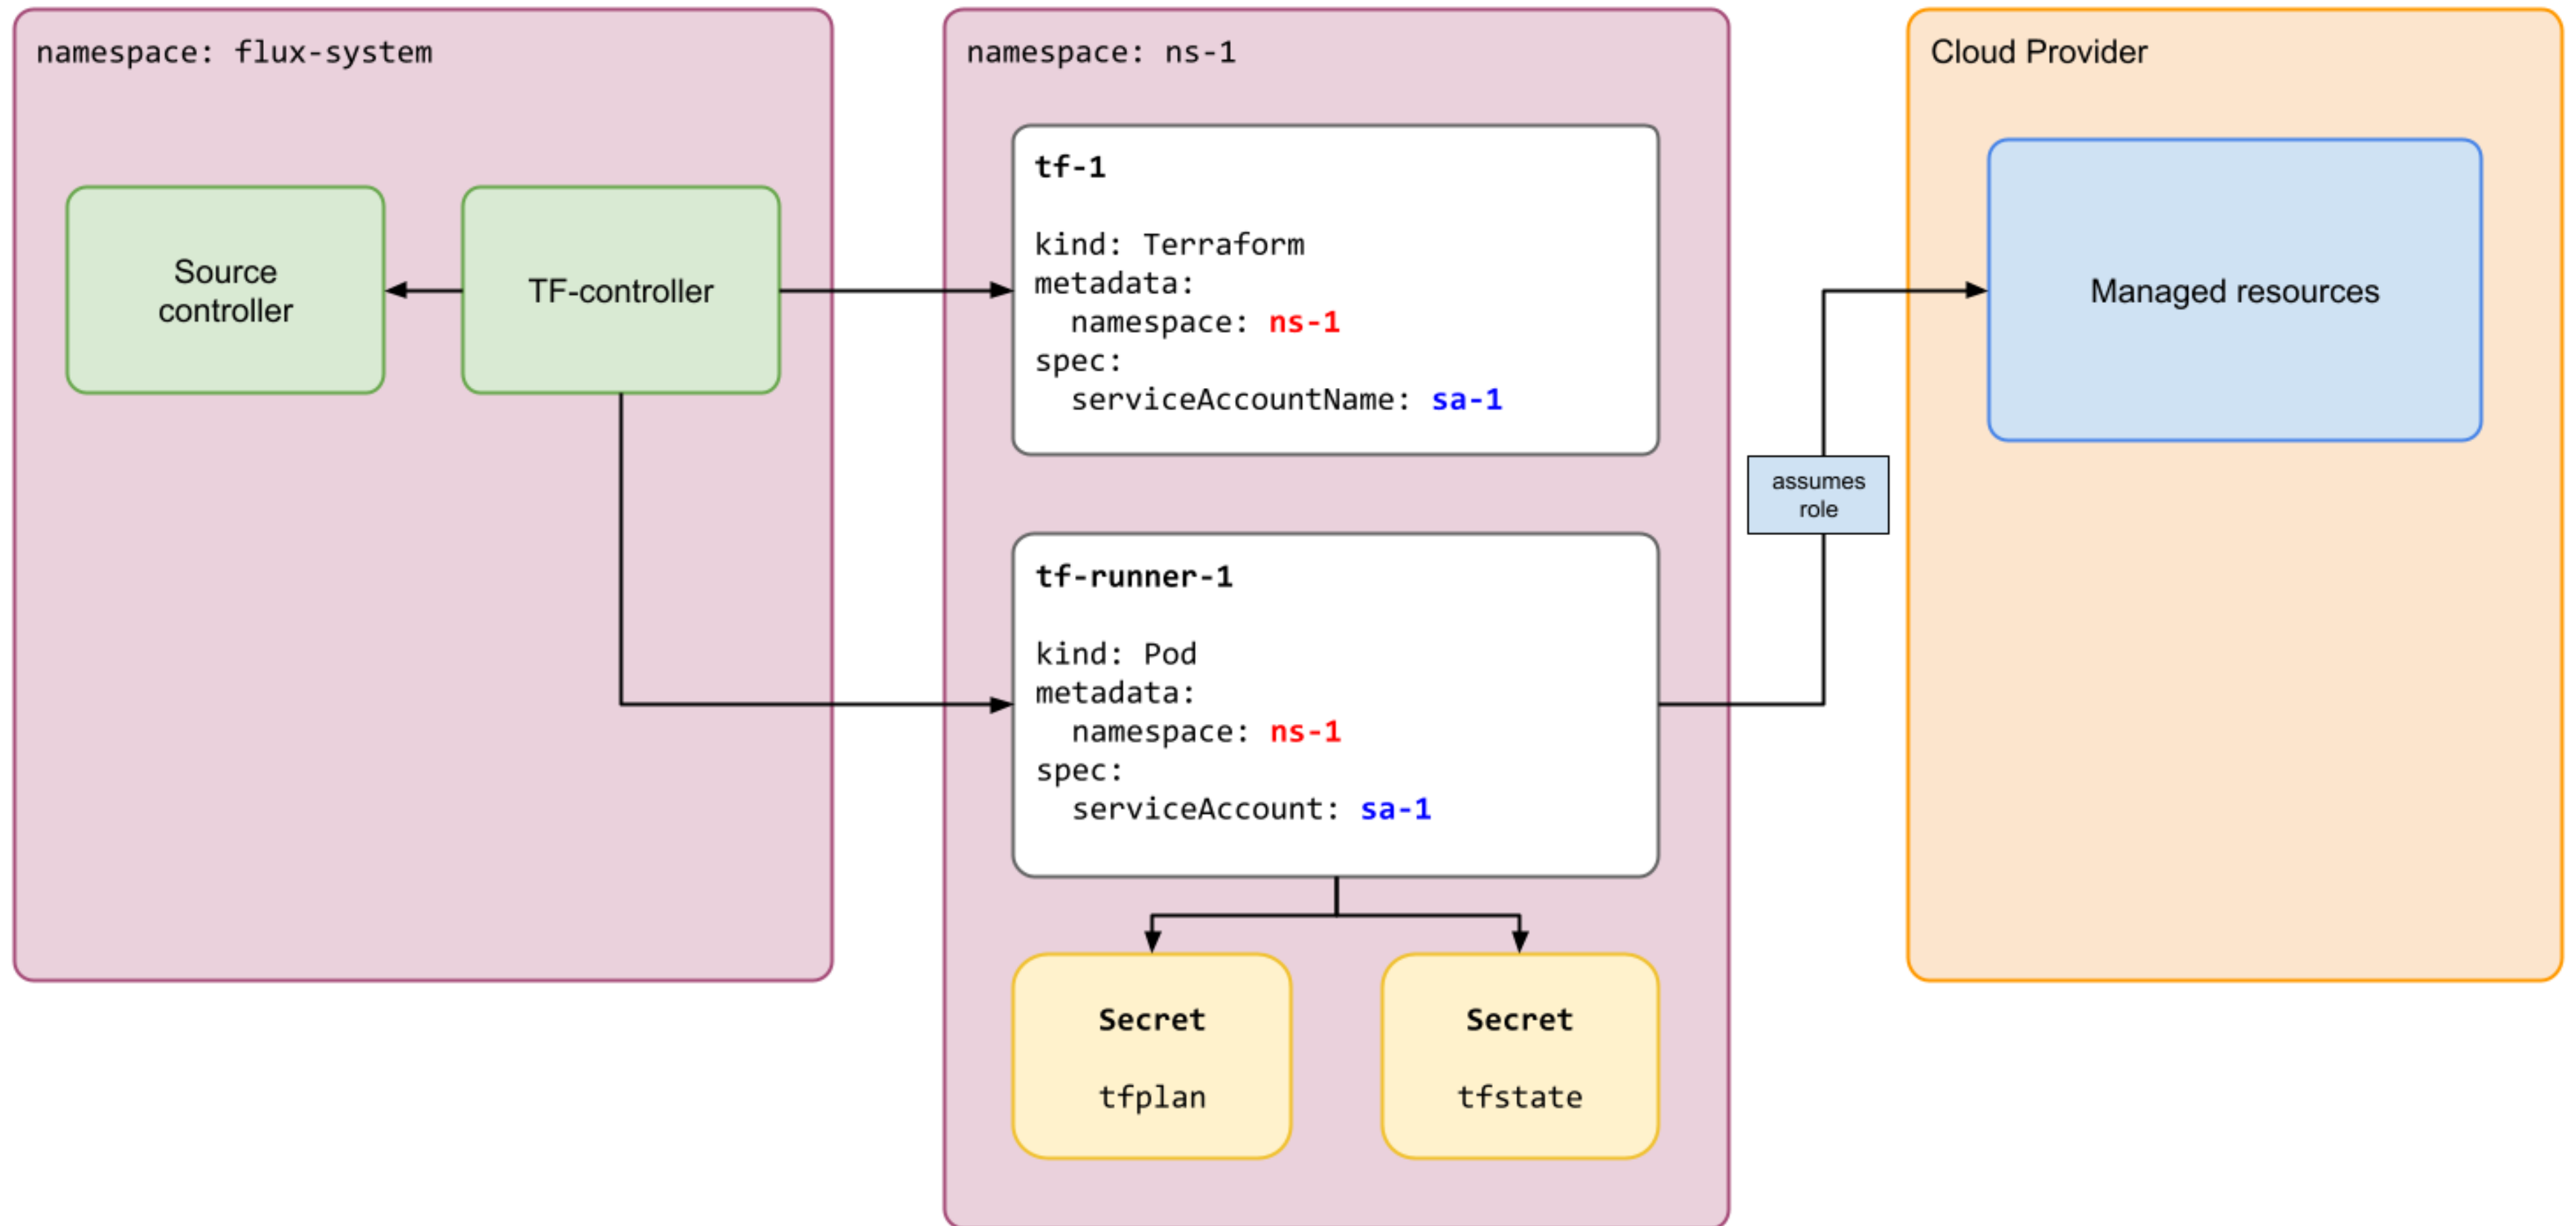
### WHAT IS GITOPS?

*GitOps is an operational framework that takes DevOps best practices used for application development such as version control, collaboration, compliance, and CI/CD, and applies them to infrastructure automation. – <https://about.gitlab.com/topics/gitops/>*

- In our context - **pull vs. push principle**
  - You don't care in *which environment* a stack runs in
  - They are ready for your stack and your code is pulled in (vs. pushed via a pipeline)



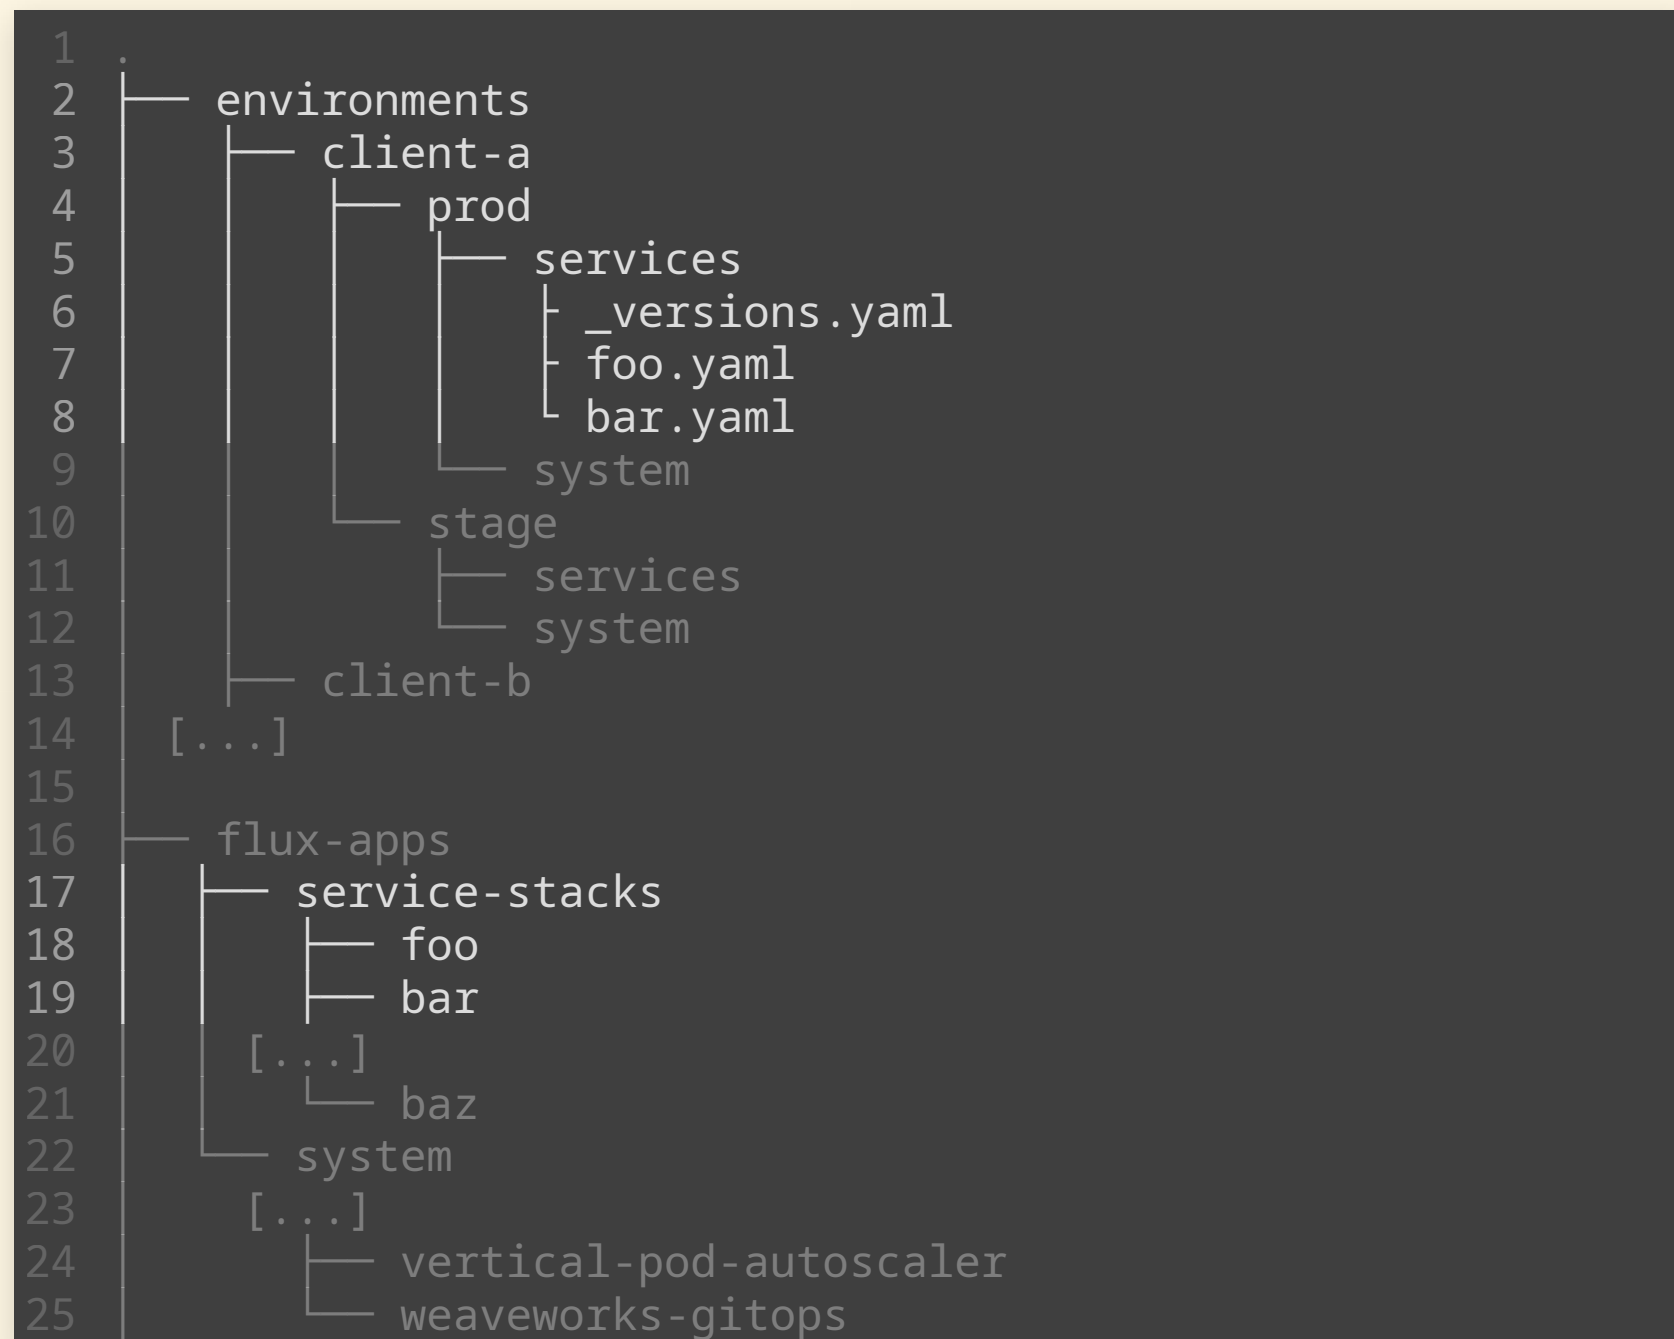
# (2.1.2) WEAVEWORKS TF CONTROLLER<sup>910</sup>



## (2.2.1) STRUCTURE OF CENTRAL FLUXCD CONFIGURATION

- Each environment must be configurable individually
  - has its own entry point for FluxCD
  - this allows for configuration of deployed services
- For audit reasons, production environments must use fixed service versions, others can use semantic versioning
- Flux applications must be DRY
  - do not c&p code
  - implication: no individual configuration of apps
- in the central Flux repo there are **NO** variables, parameters etc. pp.
  - we only document **the intent** to run a service
  - self-configuration happens inside of an environment
- Use of OCI-based registries for sources only
  - everything as a final artefact
  - `flux push artifact11` is your friend

# FLUXCD AS AN APP OF APPS SYSTEM



*from the perspective of an individual FluxCD installation*

- 0. ─ cloud and runtime is set up
  - provide data for stacks to become conscious
- 1. ─ load environment
  - primary Flux app
  - references all secondary service Flux apps
  - includes the version tracking ConfigMap
- 2. ─ load service Flux apps
  - contains relevant manifests
  - eg. OCI Sources, Terraform, Kustomization
- 3. ─ apply individual service apps



# PRIMARY FLUX APP

Applications > fluxcd-environment



SYNC  with Source  without Source

Applied Revision: main@sha1:482... Last Updated: 4 minutes ago

Applied revision: main@sha1:482acb06d5521f962959003db848ea56b533ff3

All workloads are passing health checks

> More Information

DETAILS

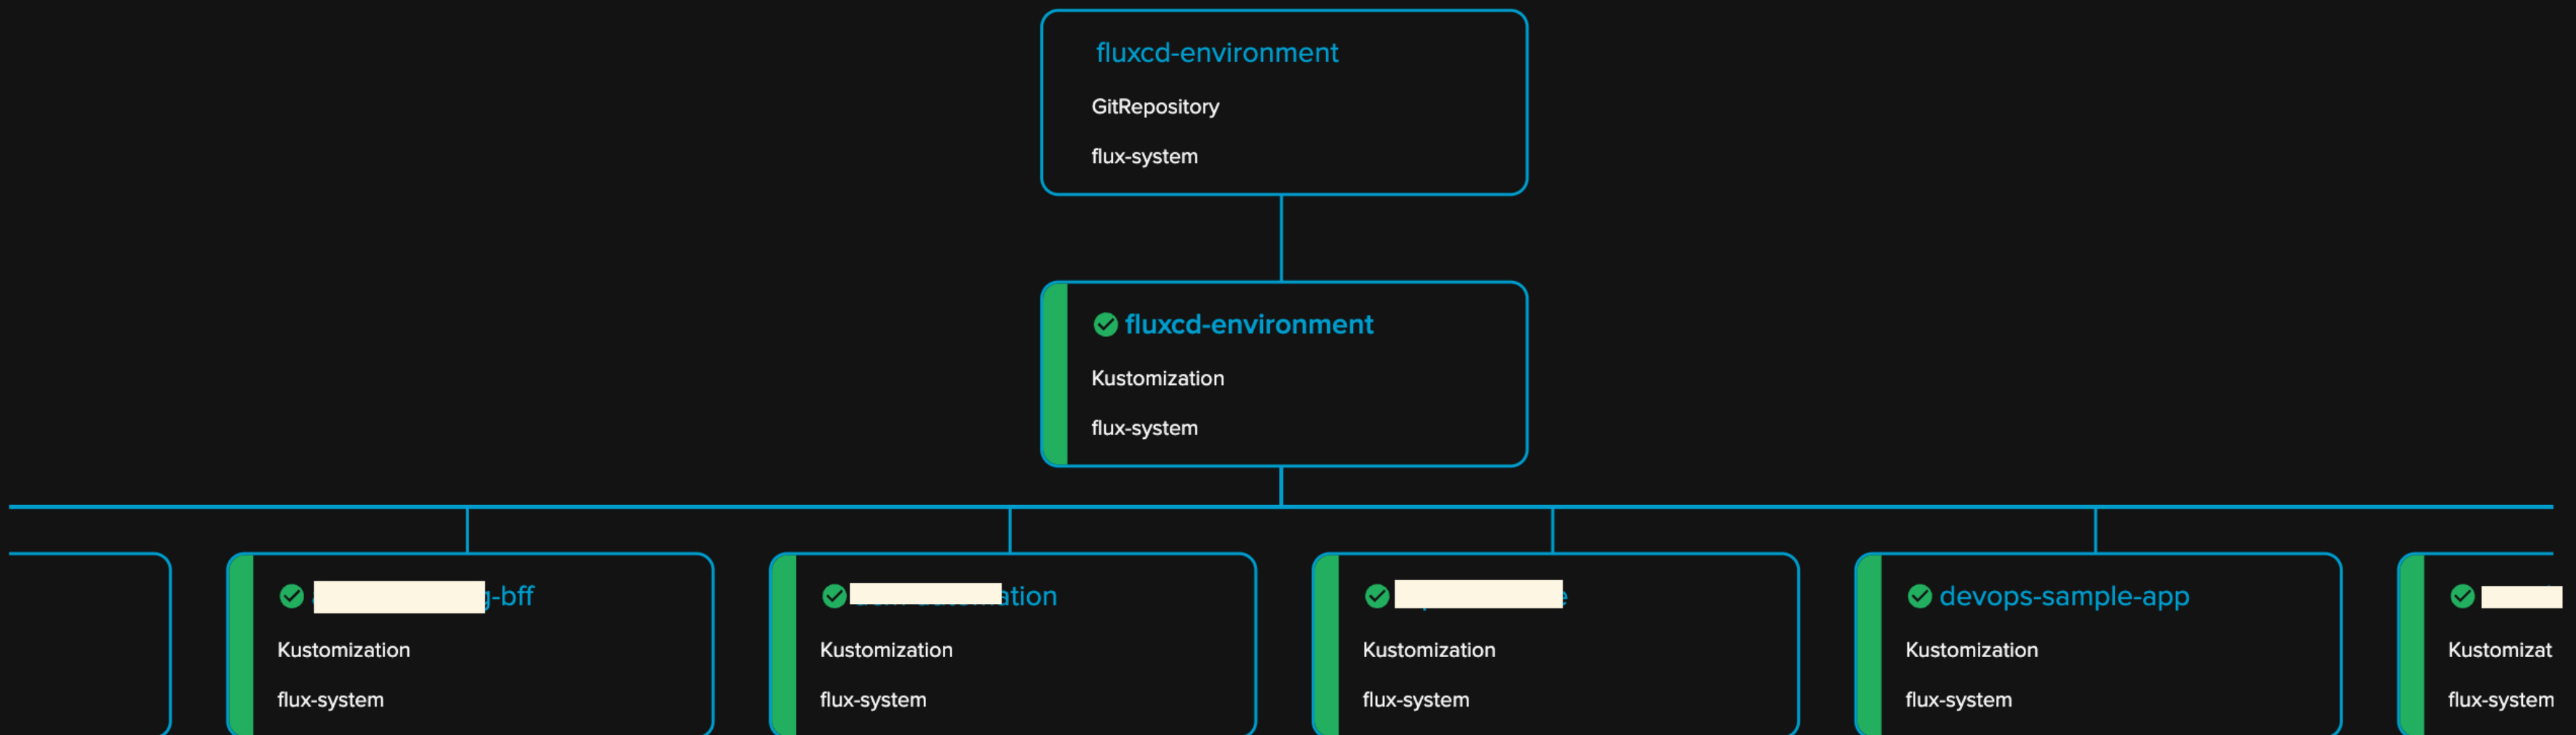
EVENTS

GRAPH

DEPENDENCIES

YAML

VIOLATIONS



# SECONDARY FLUX APP

Applications > devops-sample-app



SYNC  with Source  without Source

Applied Revision: main@sha1:482... Last Updated: 23 minutes ago

Applied revision: main@sha1:482acb06d5521f962959003db848ea56b533f1f3

All workloads are passing health checks

> More Information

DETAILS

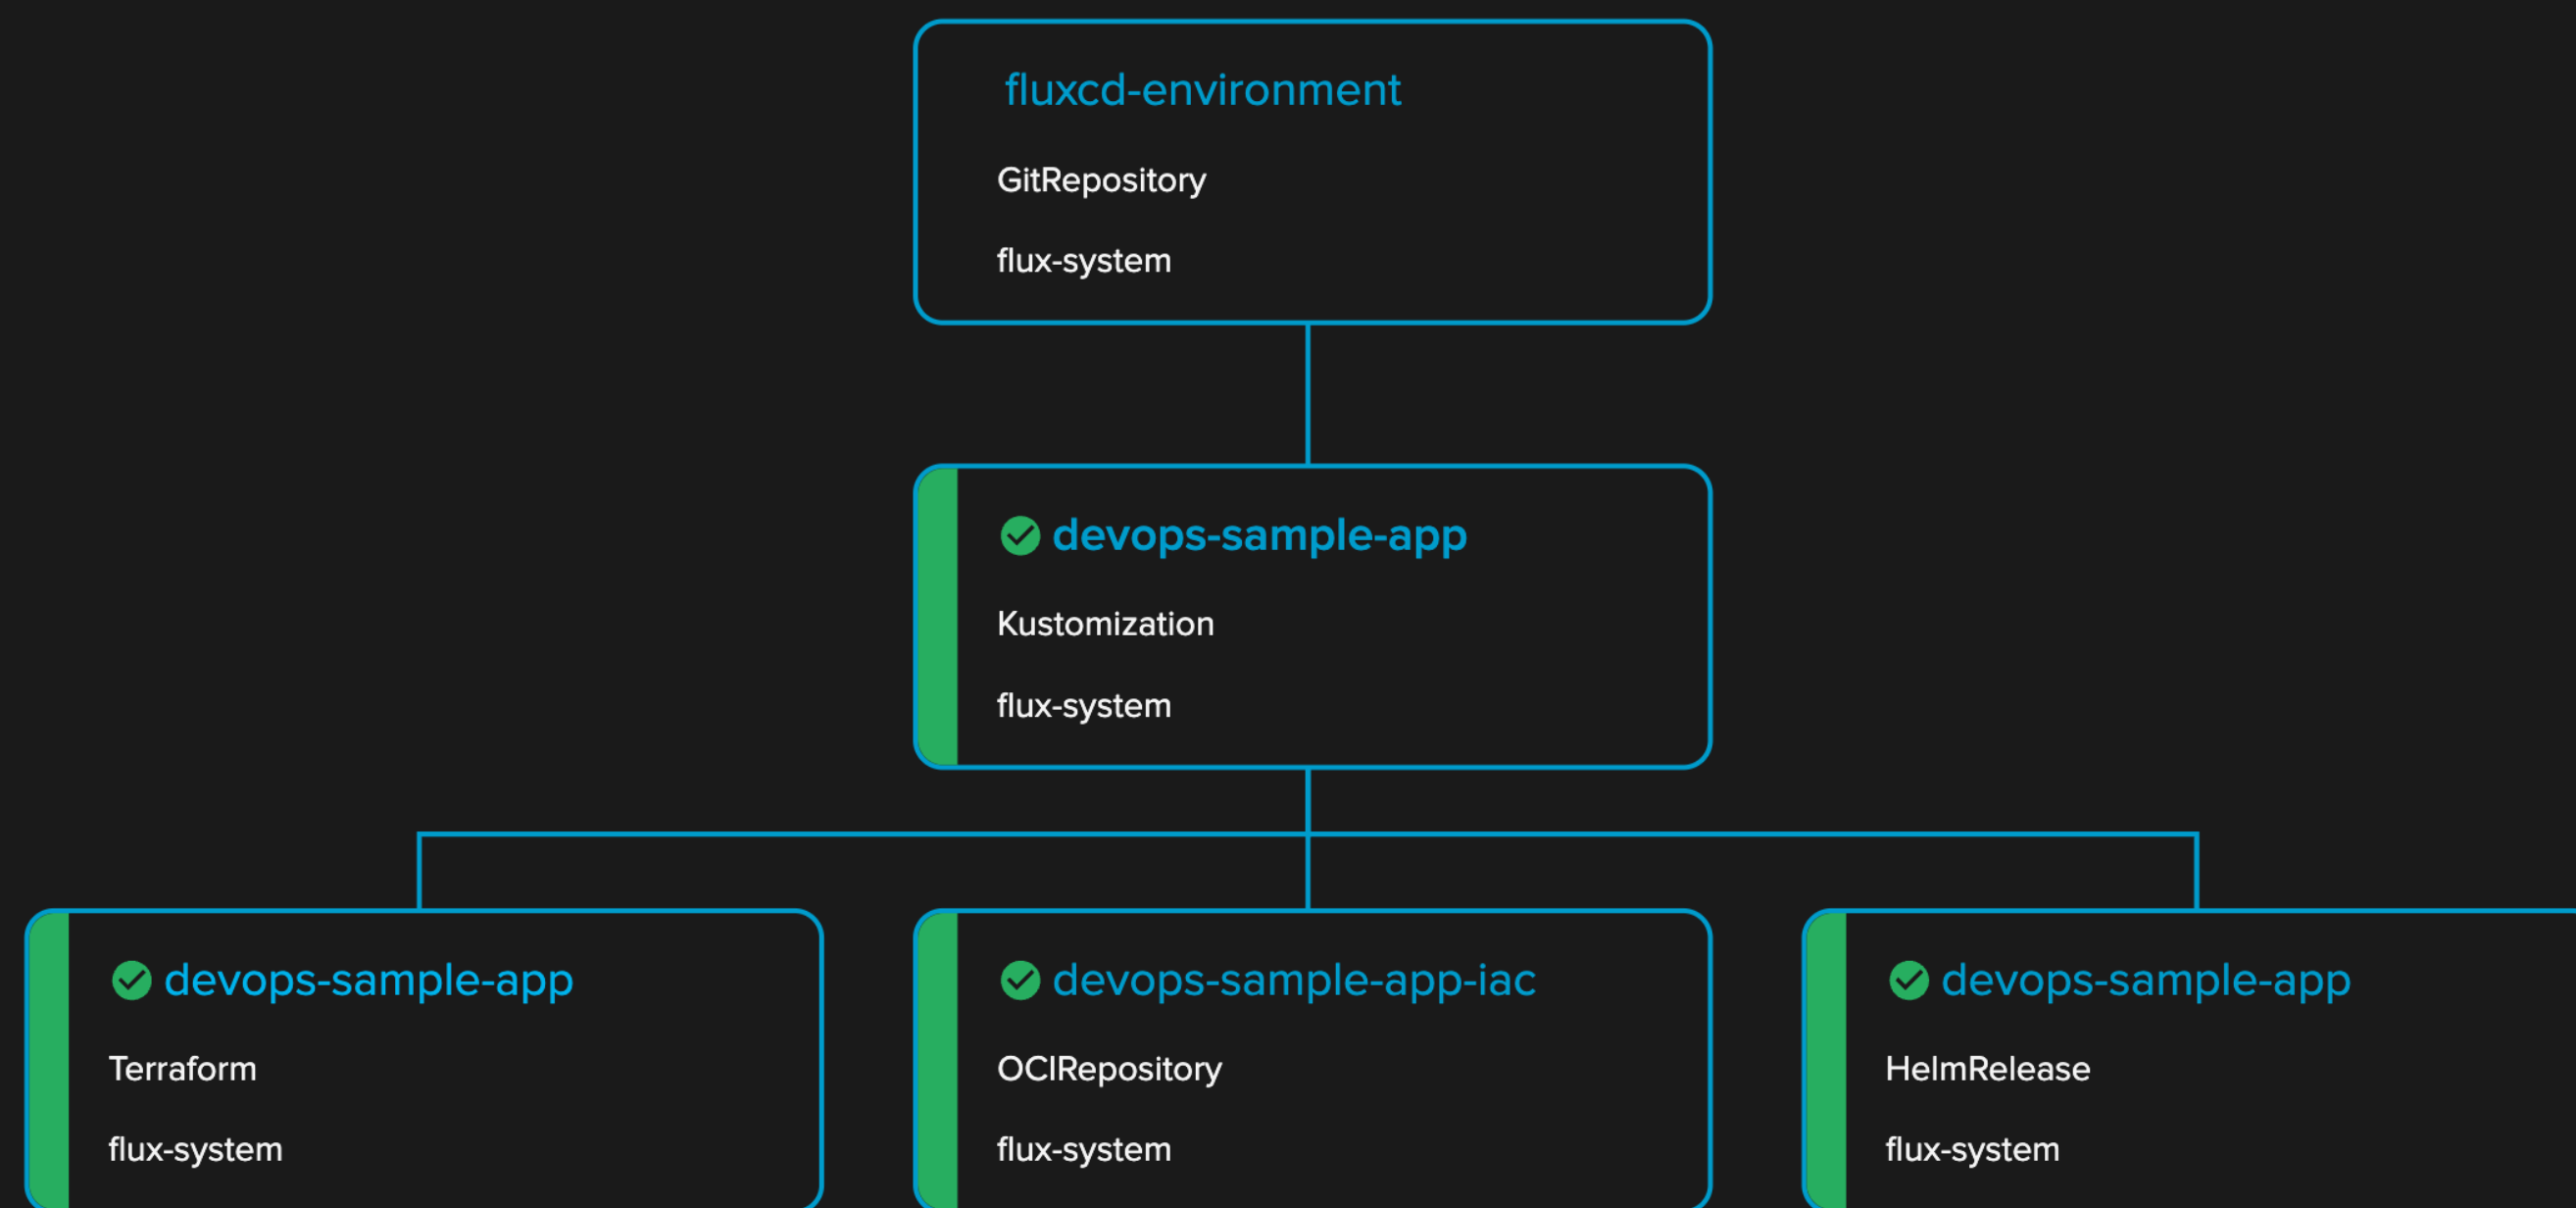
EVENTS

GRAPH

DEPENDENCIES

YAML

VIOLATIONS



85%

## (2.3.1) POST BUILD VARIABLE SUBSTITUTION<sup>12</sup>

- FluxCD's unique possibility to replace variables in rendered manifests before apply
- in FluxCD repo
  - environment specific `_versions.yaml` becomes `service-versions` ConfigMap
  - satisfies the "fixed versions" requirement
- In underlying IaC - basic environment information for a TF stack are written
  - base ConfigMap provides `client`, `environment` and other data
  - to form the path for Terraform state file

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: service-versions
5 data:
6   version_foo: "2.5.0"
7   version_foo_tf: "~ 0.1.0-0"
8   version_vertical_pod_autoscaler: "~> 9.0.0"
9   version_vertical_pod_autoscaler_tf: "~ 0.1.0"
10 ---
11 apiVersion: v1
12 kind: ConfigMap
13 metadata:
14   name: init
15 data:
16   clientId: "tenant-a"
17   domain: "stage.tenant-a.tld"
18   environment: "stage"
19   environmentClass: "non-prod"
20   region: "eu-central-1"
```

## (2.3.2) USAGE EXAMPLE

```
1 apiVersion: source.toolkit.fluxcd.io/v1beta2
2 kind: OCIRepository
3 metadata:
4   name: foo-iac
5 spec:
6   interval: 5m
7   provider: aws
8   ref:
9     semver: "${version_foo_tf}"
10  url: oci://xxx.dkr.ecr.eu-central-1.amazonaws.com/iac/foo
```

```
1 apiVersion: infra.contrib.fluxcd.io/v1alpha2
2 kind: Terraform
3 metadata:
4   name: foo
5 spec:
6   backendConfig:
7     customConfiguration: |
8       backend "s3" {
9         region      = "${region}"
10        bucket      = "terraform-states"
11        key         = "${clientId}/${environment}/stacks/"
12        role_arn    = "arn:aws:iam::xxx:role/tf-${clientId}"
13        dynamodb_table = "terraform-states-locks"
14        encrypt     = true
15      }
16   sourceRef:
17     kind: OCIRepository
18     name: foo-iac
19   vars: []
```

## (2.4.1) CONFIGURATION MANAGEMENT

- (Terraform) code is agnostic of environments
- strict division of concerns between cloud and runtime environment
  - Helm/Kustomize - Runtime (Kubernetes)
  - Terraform - Cloud
- Each **Cloud** and **Runtime environment** allow a stack to become conscious
  - Cloud: SSM data base; Runtime: ConfigMap `init`
  - 🖱️ pull of configuration vs. push
- per **code stack** - data are baked into artifact
  - terraform - single `configuration.tf`
  - kustomize - separate `overlay` directories
  - helm - separate `values.yaml`

## (2.4.2) EXAMPLE

```
locals {
  service      = "foo"
  squad       = "bar"
  domain_name = module.ssm_data.values["base"]["domain"]
  cluster_name = module.ssm_data.values["base"]["cluster"]["name"]
  client       = nonsensitive(module.ssm_data.values["base"]["client"])
  environment  = nonsensitive(module.ssm_data.values["base"]["environment"])
  env_class    = nonsensitive(module.ssm_data.values["base"]["environment_class"])

  configuration = {
    default = {
      k8s_namespace      = local.service
      k8s_sa_name        = local.service
      rds_instance_class = "db.t4g.medium"
    }
    client_a = {
      stage = {}
    }
    environment_classes = {
      non-prod = {}
      prod = {
        rds_instance_class = "db.r6g.medium"
      }
    }
  }

  # choose the right configuration based on
  # client/environment/environment class or simply defaults
  selected_configuration = merge(
    local.configuration["default"],
    try(local.configuration[local.client][local.environment], {}
  )
}
```

```
# get the central SSM config parameters
module "ssm_data" {
  source = "registry.example.com/foo/ssm_full_json_store/aws"
  version = "0.3.1"

  path = var.config_map_base_path
  include_filter_regex = "(base|foo|bar)"
}

module "database" {
  source = "registry.example.com/foo/RDS/aws"
  version = "3.5.0"

  identifier      = local.service
  squad          = local.squad
  rds_engine_version = local.selected_configuration["rds_engine_version"]
  rds_instance_class = local.selected_configuration["rds_instance_class"]
  client_id       = local.client
  environment     = local.environment
  vpc_id          = module.ssm_data.values["base"]["aws"]["vpc_id"]
  subnet_ids     = module.ssm_data.values["base"]["aws"]["subnet_ids"]
  # [...]
}
```

## (2.4.3) CONNECTING CLOUD AND RUNTIME

- remember: division of concerns - cloud and runtime
- Terraform stack writes structured data as JSON
- Runtime pulls in data via External Secrets Operator<sup>13</sup>
- Reloader watches and upgrades Pods with their associated data

```
module "ssm_service_data" {
  source = "registry.example.com/foo/ssm_json_store/aws"
  version = "1.0.2"

  path = "/configuration"
  name = "foo"
  data = {
    installed = true
    private = {
      database = {
        database_name      = module.database.databases
        database_username  = module.database.database_username
        endpoint            = module.database.endpoint
        reader_endpoint    = module.database.reader_endpoint
        port                = module.database.cluster_port
      }
    }
    public = {}
  }
}
```

```
apiVersion: external-secrets.io/v1beta1
kind: ExternalSecret
metadata:
  name: foo-secrets-ssm
spec:
  target:
    name: foo-secrets-ssm
  data:
    # [...]
    - remoteRef:
        key: /configuration/foo
        property: private.database.database_username
        secretKey: DATABASE_USER
    - remoteRef:
        key: /configuration/foo
        property: private.database.endpoint
        secretKey: DATABASE_HOST
    ---
kind: Deployment
metadata:
  annotations:
    reloader.stakater.com/auto: "true"
```

## (2.5) SPECIFICS OF TF-CONTROLLER

### (2.5.1) TRAFFIC

- each stack has its own `tf-runner` pod
  - **Decision:** no persistent pods between runs for security reasons (permissions of SA)
- Sizing example: `terraform-provider-aws_5.31.0_darwin_arm64.zip` = 84MB
- NAT costs (*AWS specific issue; GCP lowered egress costs to \$0 recently*)
  - reconcile every 30'
  - `terraform init` for each execution
  - $100 \text{ stacks} * 48 \text{ runs/day} * \sim 100\text{MB providers} * \$0,052/\text{GB} = 480\text{GB}/\$24,96 \text{ day/environment}$
- boring-registry to the rescue 🎉
  - [caching, pull-through proxy](#)
  - Provider Network Mirror Protocol
- provider stored and delivered as S3 objects
- 😊 use S3 VPC endpoints



# .terraformrc

```
credentials "my.terraform-registry.foo.bar" {  
  token = "7H151553CUr3!" # we are 1337  
}  
  
provider_installation {  
  network_mirror {  
    url = "https://my.terraform-registry.foo.bar/v1/mirror/"  
    include = ["*/*"]  
  }  
}
```

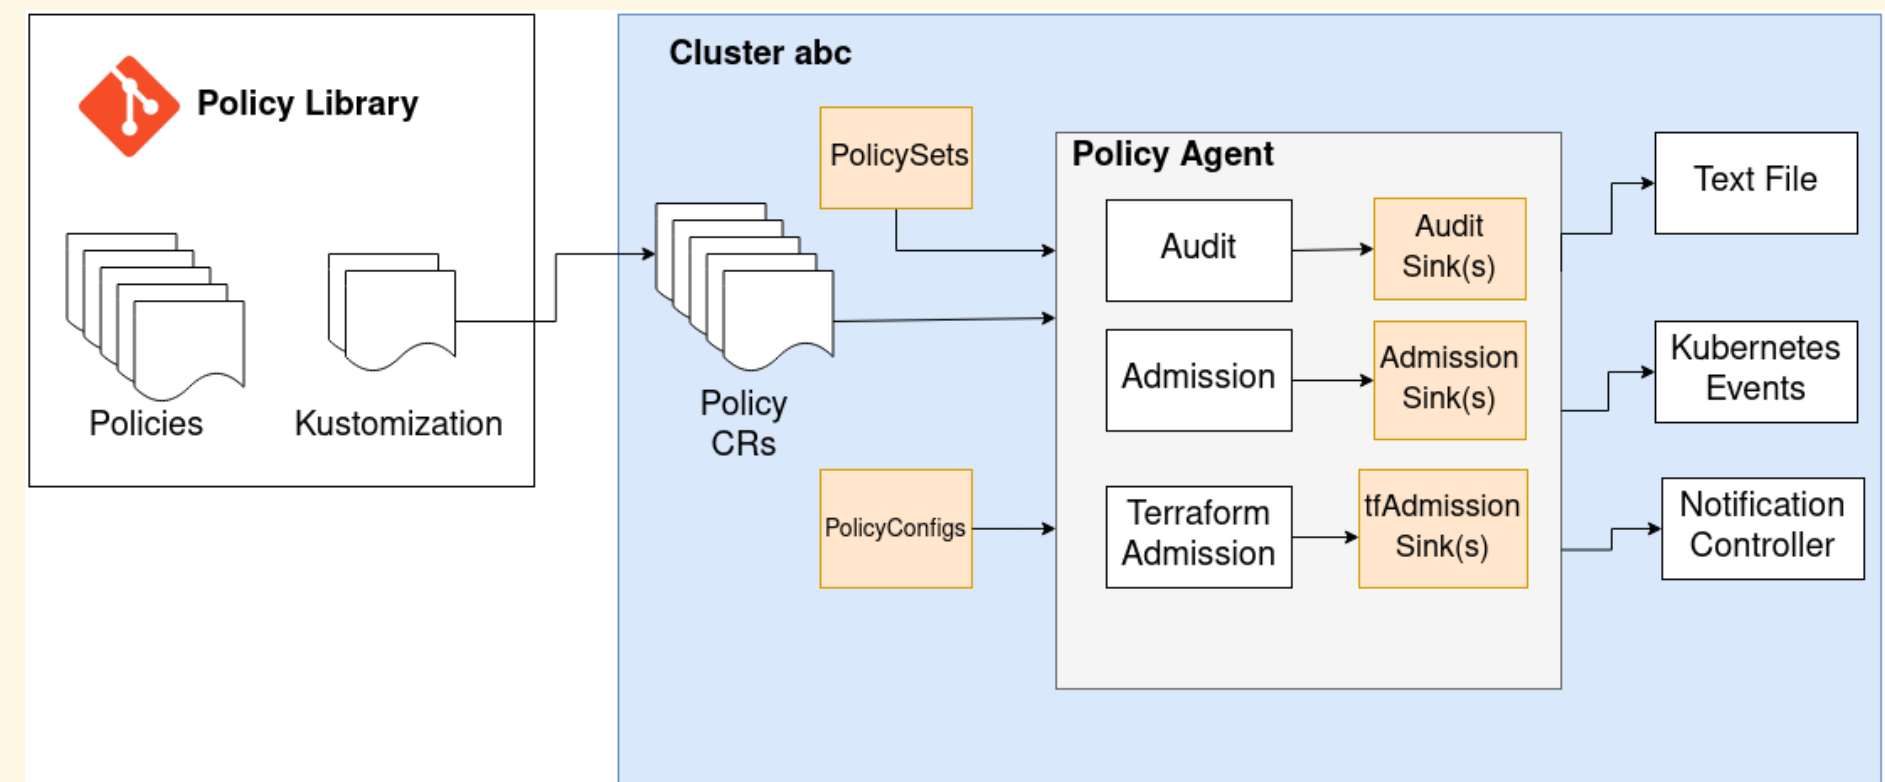
## (2.5.2) KUBERNETES RESOURCES

- each reconcile cycle triggers one `tf-runner` pod per stack
- each `tf-runner` pods consumes
  - ~800m CPU
  - ~150M Memory
- This would spawn a lot of machines at times
- using k8s limits based on `priorityClass`

```
1 apiVersion: scheduling.k8s.io/v1
2 description: used to limit the number of terraform runners
3 kind: PriorityClass
4 metadata:
5   name: terraform
6 value: 0 # same priority as everybody else
7 ---
8 apiVersion: v1
9 kind: ResourceQuota
10 metadata:
11   name: terraform-runners
12 spec:
13   hard:
14     pods: "10"
15   scopeSelector:
16     matchExpressions:
17     - operator: In
18       scopeName: PriorityClass
19     values:
20     - terraform
```

## (2.6) WEAVE POLICY ENGINE<sup>14</sup>

- based on Rego and similar to Open Policy Agent
- **Goal:** auto approve Terraform changes
  - **Decision:** no destroy/recreate
  - **Decision:** no direct IAM resources (only via controlled modules)
- ⚠ not an easy task - talk of its own



# (2.7) WEAVE GITOPS UI<sup>15</sup>

AKA - THE MISSING FLUXCD UI

**weavegitops** Applications

SYNC [dropdown] [pause] [play] [search] [filter]

| <input type="checkbox"/> | NAME                        | KIND          | NAMESPACE   | TENANT | SOURCE                                | STATUS  | ↓ | M |
|--------------------------|-----------------------------|---------------|-------------|--------|---------------------------------------|---------|---|---|
| <input type="checkbox"/> | <a href="#">flux-system</a> | Kustomization | flux-system | -      | <a href="#">flux-system</a>           | ✓ Ready |   | A |
| <input type="checkbox"/> | <a href="#">ww-gitops</a>   | HelmRelease   | flux-system | -      | <a href="#">flux-system-ww-gitops</a> | ✓ Ready |   | R |



APPLICATIONS

SOURCES

IMAGE AUTOMATION

**FLUX RUNTIME**

DOCS

This cluster is running Flux version: v0.34.0

**CONTROLLERS**

CRDS



| NAME ↓                  | STATUS  | NAMESPACE   | IMAGE                                                                                                               |
|-------------------------|---------|-------------|---------------------------------------------------------------------------------------------------------------------|
| helm-controller         | ✓ Ready | flux-system | <a href="https://ghcr.io/fluxcd/helm-controller:v0.24.0">ghcr.io/fluxcd/helm-controller:v0.24.0</a>                 |
| kustomize-controller    | ✓ Ready | flux-system | <a href="https://ghcr.io/fluxcd/kustomize-controller:v0.28.0">ghcr.io/fluxcd/kustomize-controller:v0.28.0</a>       |
| notification-controller | ✓ Ready | flux-system | <a href="https://ghcr.io/fluxcd/notification-controller:v0.26.0">ghcr.io/fluxcd/notification-controller:v0.26.0</a> |
| source-controller       | ✓ Ready | flux-system | <a href="https://ghcr.io/fluxcd/source-controller:v0.29.0">ghcr.io/fluxcd/source-controller:v0.29.0</a>             |

# (3.0) IS IT PRODUCTION READY?

- tf-controller is sometimes uncertain about the state
- slow development of tf-controller, ~~thank you~~ HashiCorp
  - in principle ready for OpenTofu<sup>16</sup>
  - the talk uses features from a pre-release<sup>17</sup>
- observability is not ideal
  - eg. finding all Terraform Manifests, which have a pending plan

## Be honest, where are you in the project?

- In the middle of cutting the large TF stacks
  - 🖱️ very useful tool: [minamijoyo/tfmigrate](https://github.com/minamijoyo/tfmigrate)
- Automatic approvals are yet to come
- **Branch Planner** needs to be implemented to enable full developer ownership
- after IaC migration, services move to FluxCD as well

## (3.1) WHY NOT THE BACK STACK<sup>18</sup>?

- **Backstage (B)**: A self-service portal to empower developers
- **Argo CD (A)**: A GitOps-based continuous delivery (CD) tool for streamlined software delivery.
- **Crossplane (C)**: A universal control plane simplifying self-service infrastructure provisioning through abstractions.
- **Kyverno (K)**: A Policy as Code (PaC) tool
- existent Terraform stack and knowledge did not justify re-write of IaC
- Crossplane is bound to 1 kubernetes cluster (state in etcd) where Terraform is bound to a state file
- Introduction of Backstage was out of scope
- ArgoCD vs. FluxCD
  - ArgoCD's handling of Helm charts (templated and applied)
  - TF-Controller as part of FluxCD eco-system
  - ArgoCD has UI and concept of multi-cluster baked in
- Kyverno “runs as a dynamic admission controller” can not be used as a decision engine

## (3.2) DOWNSIDES

- development and local testing of TF code is hard
  - possibly via Branch Planner
  - *only for Github sources*
- Terraform module registry - so batteries included for developers? yes, kind of, but
  - Terraform understanding needed
  - it is hard to grock the stack data exchange concept
  - we provide template repositories, use case documentation
- TF-Controller: (un)interruptable pods needed (for writing states)
- missing UI (for TF-Controller) and Monitoring APIs
- implicit data contracts between Terraform stacks



# (3.2.1) - AN UNCERTAIN FUTURE

The image shows a LinkedIn post from Alexis Richardson, CEO of Weaveworks, dated 2 days ago. The post text reads: "Hi everyone I am very sad to announce - officially - that Weaveworks will be closing its doors and shutting down commercial operations. Customers and partners will be working with a financial trustee whom we shall announce soon." The post is overlaid on a TechCrunch article titled "Cloud native container management platform Weaveworks shuts its doors" and a The Register article titled "GitOps pioneer Weaveworks unravels after funding fabric frays".

**LinkedIn Post:**

- Author: Alexis Richardson, CEO Weaveworks
- Text: "Hi everyone I am very sad to announce - officially - that Weaveworks will be closing its doors and shutting down commercial operations. Customers and partners will be working with a financial trustee whom we shall announce soon."

**TechCrunch Article:**

- Title: Cloud native container management platform Weaveworks shuts its doors
- Category: Startups
- Date: April 25, 2024
- Location: Boston, MA
- Call to Action: Register Now

**The Register Article:**

- Title: GitOps pioneer Weaveworks unravels after funding fabric frays
- Subtitle: Company burned through \$61.6M in investment
- Category: CXO
- Comments: 3

<https://www.linkedin.com/feed/update/urn:li:activity:7160295096825860096/>

## (3.3) UPSIDES

- all domains (code, kubernetes and cloud environment) follow the same pattern
  - same CI and CD
  - same artefact type (OCI)
  - similar release cycles
  - single entry point for Product Owners
- IaC runner can be replaced
  - TF-Controller is just a controlled terraform executor
  - migration to eg. Spacelift.io or others possible
  - break-the-glas scenario supported (manual stack execution)
- Terraform/OpenTofu eco-system can be reused
  - providers
  - knowledge and modules

# (3.4) THANKSIDES

- LYNQTECH GmbH for granting permission to share information and code
  - 😊 LYNQTECH is hiring <https://www.lynq.tech/jobs/>
- All colleagues who were and are part of this journey
- The FluxCD Community and WeaveWorks for their software



LinkedIn



Website



GitHub

# (4.0) ARCHITECTURAL DECISIONS

## General FluxCD

- Each tenant environment must be configurable individually
  - For audit reasons, production envs must use fixed service versions
- Applications, in the central repo, must be DRY. No individual stacks.
- Use of OCI-based registries for sources only (exception: external Helm)
- Code is agnostic of environments and is not parameterised
  - Each **cloud (AWS)** and **runtime (Kubernetes) environment** allows a stack to become concious
  - kustomize style data baked into artifact
- Secrets synchronised via External-Secret Operator
- Kubernetes cluster should be treated as cattle

## TF-Controller

- No vendor lock-in; re-usability of eco-system strong plus
  - Terraform providers
  - Terraform OSS modules
- No persistent pods between runs
- Aim for Auto approval for Terraform changes
  - no destroy/recreate
  - no direct IAM resources (only via controlled modules)
  - only approved top-level module sources

# SOURCES + LINKS

1. FluxCD documentation - <https://fluxcd.io/flux/components/>
2. Weave GitOps // Terraform Controller documentation - <https://weaveworks.github.io/tf-controller/>
3. Weave GitOps // The Policy Ecosystem - <https://docs.gitops.weave.works/docs/policy/getting-started/>



1. your experience might be different 😊↔
2. <https://developer.hashicorp.com/terraform/language/files#the-root-module>↔
3. [https://en.wikipedia.org/wiki/Don%27t\\_repeat\\_yourself](https://en.wikipedia.org/wiki/Don%27t_repeat_yourself)↔
4. HashiTalks DACH 2020 - [Opinionated terraform modules and a registry](#)↔
5. [How TIER switched paradigms - from team- to service-centric](#)↔
6. [Sensitive Data in State](#)↔
7. [TF-CIX as an approach to share information between terraform stacks](#)↔
8. <https://fluxcd.io/flux/components/>↔
9. <https://github.com/weaveworks/tf-controller>↔
10. ***Please note:*** *As the tf-runner ServiceAccount is usually very powerful, do not run it in an accessible namespace!*↔
11. [https://fluxcd.io/flux/cmd/flux\\_push\\_artifact/](https://fluxcd.io/flux/cmd/flux_push_artifact/)↔
12. <https://fluxcd.io/flux/components/kustomize/kustomizations/#post-build-variable-substitution/>↔
13. <https://external-secrets.io/>, [stakater/Reloader](#)↔
14. [Weave Policy Engine, Integrate TF Controller with Flux Receivers and Alerts, Open Policy Agent](#)↔
15. <https://github.com/weaveworks/weave-gitops> and <https://docs.gitops.weave.works/>↔
16. <https://www.opentofu.org/>↔
17. <https://github.com/weaveworks/tf-controller/releases/tag/v0.16.0-rc.3>↔